



Using AI to Predict Bangladesh's Drinking Problem

Kane Swartz

941226

A dissertation submitted in partial fulfilment
of the requirements for the degree of
Bachelor of Science
of the
University of Portsmouth.

School of Computing
Research Project

2023

Declaration

No portion of the work contained in this document has been submitted in support of an application for a degree or qualification of this or any other university or other institution of learning. All verbatim extracts have been distinguished by quotation marks, and all sources of information have been specifically acknowledged.

Date: 2023

Abstract

The drinking problem referred to in the title is groundwater arsenic pollution. This project aims to evaluate existing expert system predictive models developed in the University of Portsmouth and compare them to machine learning based models.

In 2023 machine learning and artificial intelligence are becoming ever more prevalent, this dissertation explores the potential impact this field could have on predicting the arsenic pollution of groundwater.

Acknowledgements

iArsenic has been foundational to this project. The dataset extracted from iArsenic has been a huge asset and represents hours of high quality work from all of the iArsenic team who have ensured source data has been cleaned and aggregated to a high standard.

Thank you to Dr Mohammed Hoque, the initial visionary for iArsenic and client to this project. Dr Jacek Kopecký for leading the development of the iArsenic web application and the following contributors in the iArsenic team:

- Dillon O'Shea (TingeOGinge)
- Lewis Knewton (lewisknewton)
- Nylah Daniella Klasson (TxJson)
- Benjamin Kyd (benkyd)
- Fionn Plant (CybrChick007)
- AloysT
- PreslavPO

Additional thanks to my supervisor Dr Rich Boakes and friend and mentor Dr David Ralph, who have each provided key direction leading up to and throughout the development of this project.

Consent to Share

I consent for this project to be archived by the University Library and potentially used as an example project for future students.

Contents

1	Introduction	10
1.1	Background and Historical Context	10
1.1.1	Solution Approach	10
1.2	What is iArsenic	11
1.3	Research Questions	11
1.4	Conclusion	11
2	Literature Review	12
2.1	Existing Work	12
2.1.1	iArsenic	12
2.1.2	Predicting Arsenic Concentrations with Surface Parameters . . .	12
2.1.3	Satellite Mapping of Flooding Patterns	13
2.2	The Black Box Problem	13
2.3	Tool Selection	14
2.3.1	Model Implementation	14
2.3.2	Python Dependency Management	15
2.4	Model Type	16
2.4.1	Validating Models & Minimizing Dataset Bias	17
2.4.2	Evaluation Methods	18
2.5	Conclusion	20
3	Methodology	21
3.1	Research Potential Projects	21
3.2	Continuously Identifying Client Requirements	22
3.2.1	Clarifying the Project Objectives	22
3.3	Research Applied Machine Learning Methodologies	23
3.4	Understanding the iArsenic Code	23
3.4.1	Producing a standard dataset	23
3.4.2	Generating iArsenic models from a dataset	23
3.4.3	Generating predictions from generated iArsenic models	24

3.5	Developing New Models	24
3.6	Integrating iArsenic and scikit-learn	27
3.7	Main Project Structure	27
3.8	Conclusion	28
4	Experiment Design	29
4.1	Well Data in iArsenic	29
4.1.1	Geodata	35
4.2	iArsenic Integration	35
4.2.1	Data Preparation	35
4.2.2	Generating the iArsenic Models	35
4.2.3	Interfacing With Generated iArsenic Models	36
4.2.4	Imputing Stain Colour	36
4.3	Creating a Common Model Interface	40
4.3.1	Implementing the Common Model Interface	40
4.4	Evaluating the Models	40
4.4.1	Standardizing Classifications	41
4.5	Creating a Main Script to Run & Evaluate Experiments	41
4.5.1	Main Script Structure	41
4.6	Development of New Models	42
4.6.1	Summary of Models and Preprocessing Techniques Implemented	42
4.6.2	model6	43
4.6.3	model7	43
4.6.4	model8	43
4.6.5	model9	44
4.6.6	model10	45
4.6.7	model11	45
4.6.8	model12	45
4.6.9	model13	46
4.6.10	Model14	46
4.7	Conclusion	46
5	Results and Evaluation	50
5.1	Verifying Bias in Test Train Split	50
5.1.1	Quantifying Bias	50
5.2	Desired Outcome 1: Evaluating iArsenic Models	51
5.2.1	Benefit of Evaluation of iArsenic Models	51
5.2.2	Desired Outcome Achievement	52
5.3	Desired Outcome 2: Comparing iArsenic to Machine Learning Models . .	53

5.3.1	iArsenic & Machine Learning Models Performance	53
5.3.2	Desired Outcome Achievement	53
5.4	Conclusion	53
6	Limitations & Future Work	56
6.1	Limitations	56
6.2	Future Work	57
6.2.1	Further Data Integration	57
6.2.2	Different Model Types	59
6.2.3	Further Development of iArsenic Models	60
6.3	Conclusion	61
7	Reflection & Conclusion	62
7.1	Research Questions	62
7.1.1	What is the performance of the existing iArsenic models	62
7.1.2	Do machine learning models have potential in the prediction of groundwater arsenic pollution?	62
7.2	New Models Development Strategy	63
7.3	Achievement of Personal Goals	63
7.4	Conclusion	63
	References	64
A	Appendix	66
A.1	Colorblind Compatible Figures	66
A.2	PID	66

List of Tables

4.1	Region sizes, datapoints per region & count	31
4.2	Model8 MLP configuration	44
4.3	Model9 MLP configuration	44
4.4	Model14 MLP configuration	46
5.1	Standard deviation of specificity and sensitivity across each test train split	51
5.2	Table of evaluation metrics for all models	55

List of Figures

3.1	Package.json snippet showing script to load data from iArsenic	23
3.2	Flowchart showing process to implement model in iArsenic webapp . . .	24
3.3	Python code used to programmatically generate iArsenic models. In iArsenic this is done via the command line	25
3.4	Code used in iArsenic to produce estimate from a model in the webapp . .	25
3.5	Summary of results of supervised model implementation practice	26
3.6	Sequence diagram showing Python to NodeJS variable bridge over storage	27
3.7	Project file structure	28
4.1	Distribution of Well Data by District in iArsenic Standard Dataset	30
4.2	The Divisions of Bangladesh mapped and labelled	32
4.3	Average (mean) Well Depth by Upazila	33
4.4	Chart showing the relationship between mean well depth and arsenic pollution	33
4.5	Distribution of Average (mean) Arsenic by Upazila in the iArsenic Standard Dataset	34
4.6	iArsenic model performance by accuracy for default red vs black staining	37
4.7	Model sensitivity and specificity stain parameter 'Black' (top) vs 'Red' (bottom)	38
4.8	Confusion Matrix model5 'Black' (top) vs 'Red' (bottom) Staining	39
4.9	Snippet showing m8 optimization method, where a different model is trained for each of the 8 Divisions	48
4.10	Distribution of Well Data by District in iArsenic Processed Dataset	49
5.1	Model performance by f1 score	54
6.1	Model9 loss curve 50, 2 hidden layers	58
6.2	Model9 loss curve 2000, 1000, 100 hidden layers	58
A.1	Red excluded for red / green colourblind: Distribution of Average (mean) Arsenic by Upazila in the iArsenic Standard Dataset	79

Chapter 1

Introduction

1.1 Background and Historical Context

Arsenic pollution in groundwater is a problem faced worldwide (A. Khan, 2014). South-east Asian populations, particularly those of Bangladesh, are arguably the most negatively impacted by arsenic contamination in drinking water worldwide.

(Smith et al., 2000).

These tube wells were typically installed at a depth of under 200 meters. Arsenic concentrations in groundwater peak at between 15 and 25 meters of depth and are typically only below the World Health Organisation's safe arsenic limit of $10\mu\text{g/litre}$ (Organization, 2003) at 200 meters of depth or more (Chakraborti et al., 2010).

Prior to the discovery of harmful levels of arsenic in the water from these tube wells, their installation was acclaimed as a huge success with regards to providing the population with clean drinking water (A. Khan, 2014). Tubewells today are still extremely popular in Bangladesh with approximately 95% of rural people and 70% of urban people using them in 2020 (Ghosh et al., 2020).

1.1.1 Solution Approach

Smith et al., 2000 emphasizes that the core of any solution must be to provide non-polluted water.

This article makes it clear that there is no single solution to this issue, it may seem simple to provide water filters, for example, but for their successful deployment, they must be provided with education and community engagement to ensure their efficacy. They are also not a permanent solution.

N. I. Khan and Yang, 2014 illustrates the scale and complexity of a comprehensive solution approach by analysing the motivations of key stakeholders.

1.2 What is iArsenic

iArsenic is a web application that estimates the arsenic concentration for a well in Bangladesh based on its geographic region, depth and visible staining on the well.

This is achieved by aggregating source data about these wells and using that aggregate data to produce expert system prediction models which are available via a web application.

The tooling developed to produce these models are general purpose and modular, enabling them to be imported into other projects. The time and expertise required to clean and aggregate the source data used by these models are very substantial and leveraging this existing work allows higher level outcomes to be brought into scope.

The data in iArsenic has been selected from reputable sources and represents a large and high quality dataset, representing the majority of Bangladesh.

1.3 Research Questions

1. What is the performance of the existing iArsenic models.
2. Do machine learning models have potential in the prediction of groundwater arsenic pollution?

1.4 Conclusion

Arsenic pollution is a serious and ongoing issue and it is a great shame that despite relevant organisations being aware of this problem since at least the year 2000 it persists to today in 2023.

This project does not aim to conquer and eliminate this problem, but does aim to make a contribution to the ongoing research into the use of predictive models to predict groundwater arsenic pollution.

To achieve this contribution, this project aims to validate existing predictive models and provide an example for the application of machine learning in this field.

Chapter 2

Literature Review

This chapter outlines the existing systems being used to predict the safety of drinking water sources in Bangladesh and other regions where applicable, the tooling and techniques required to implement and evaluate new prediction models and touches on the considerations and limitations of these models with regard to the black box problem.

2.1 Existing Work

2.1.1 iArsenic

The existing iArsenic models are expert system models. This type of model follows rules defined by an expert in the field, this can be thought of as a flow chart. This approach has the benefit of the output of the model being explainable and justifiable; a white box model.

Expert systems can be limited compared to machine learning-based models as the decision boundaries made by an expert system are defined by an expert, whereas machine learning algorithms will statistically identify decision boundaries.

In total there are 5 iArsenic models, each has been developed sequentially, improving on the one before. Initially written in R, 3 of the iArsenic models, model3, model4 and model5, have been converted to NodeJS / JavaScript and uploaded to GitHub.

The iArsenic models are documented at this link: <https://github.com/portsoc/iArsenic/tree/master/preprocessing>

2.1.2 Predicting Arsenic Concentrations with Surface Parameters

Winkel et al., 2008 page 536 used latitude, longitude and surface parameters from the Food and Agriculture Organization of the United Nations to predict arsenic concentration in groundwater. This was done by calculating the correlation coefficient between these surface parameters and arsenic concentrations for use as global parameters in a classification model. This model achieved a true positive and true negative rate of 65% over

multiple countries in Southeast Asia.

2.1.3 Satellite Mapping of Flooding Patterns

Connolly et al., 2022 page 928 uses high-resolution satellite data combined with surface parameters to predict groundwater arsenic concentrations within an area of 30m using a random forest model. The results in this article are not comprehensive, though a figure for a sensitivity rate of 95% for wells with arsenic concentrations of over 10 μ g/l is provided.

The compromise for this high sensitivity level appears to be that the dataset covers a smaller area than the iArsenic dataset, which covers the majority of Bangladesh or Winkel et al., 2008's dataset which covers several countries in Southeast Asia.

Conclusion

While Connolly et al., 2022's model achieving a sensitivity rate of 95% is the highest for any model found in this literature review, this does not provide a comprehensive performance evaluation for the model. A model which always predicts true will have a sensitivity of 1.

Winkel et al., 2008's model provides a more robust evaluation by providing a receiver operator curve.

It is not possible to generate a receiver operator curve for the iArsenic models. This is because the models do not provide a feature to change the classification threshold. It is possible to generate the f1 score, sensitivity and specificity. This will factor in the true positive rate and true negative rate. See these terms defined at 2.4.2 on page 18.

2.2 The Black Box Problem

Fleming et al., 2021 states that acceptance of machine learning-based approaches have faced resistance in earth & environmental sciences due to the inability to explain how machine learning models are working and the lack of understanding of artificial intelligence in the earth & environmental field. While this is partially true in terms of the black box problem there are examples of machine learning being used to analyze groundwater quality. A more difficult question to answer is whether these models are less applicable in real life compared to human-designed models, as is asserted by Guidotti et al., 2018.

In machine learning, the black box problem is based on the inability to explain how a model produces an output for a given input.

Castelvecchi, 2016 illustrates this issue with the following, imagine a black box machine

learning model was created with a set of images of mammograms including labels indicating whether the person in the scan went on to develop breast cancer. It could happen that this model outperforms human medical professionals in predicting whether someone would go on to develop breast cancer.

One could argue that the benefit of this model would be that it could identify when women should have a preventative mastectomy. However, critics assert that critical decisions cannot be made without justification and that one cannot justify a decision that cannot be explained (Loyola-Gonzalez, 2019). The human doctor's opinion will always be explainable, by the doctor themselves.

Not all machine learning base models are considered to be black box systems. Decision tree models for example are powerful machine learning tools which produce models considered white box systems meaning you can analyse them to determine how they're working (Caruana, 2006). Decision tree models can be contrasted with neural network style models which are almost always considered black box models.

Because determining whether a drinking water source is safe does factor into critical decision-making, it is important to explore the black box problem. The purpose of this project, however, is to compare the performance of a predictive model developed with machine learning to existing expert system models. Therefore whether a model is black or white box will not be a critical factor in model selection.

2.3 Tool Selection

2.3.1 Model Implementation

Preliminary investigations into machine learning, including speaking to local experts at the University of Portsmouth, identified three potential systems for designing and implementing machine learning frameworks, scikit-learn, WEKA and TensorFlow.

scikit-learn

Scikit-learn provides a well defined Python interface for a wide range of well defined machine learning models and techniques. This makes scikit-learn good for implementing existing model architecture but limits the development of new architectures. This makes scikit-learn less suitable for custom architectures and deep learning than tools such as TensorFlow.

Its Python interface does mean however that it does not limit the possibilities of pre-processing or post-processing the input and output of a model (scikit-learn, 2023).

TensorFlow

TensorFlow contrasts with scikit-learn by allowing full access to the model architecture.

This makes it suitable for both deep learning and supervised learning. TensorFlow can also follow a distributed architecture, allowing it to scale to larger datasets. This versatility however greatly increases the complexity of the tool in comparison to WEKA or scikit-learn (Google, 2023).

TensorFlow provides GPU support, which could be used to decrease the time taken when testing and training, allowing less optimization to be applied to models.

WEKA

Of the three tools for model implementation considered, WEKA is the most simple to use, providing a graphical user interface. While this simplicity means WEKA is very easy to use, it is also limited in its versatility, providing a limited API, making it difficult to integrate with the iArsenic models and fewer model types and configuration options compared to scikit-learn or TensorFlow (WaikatoUniversity, 2023).

Conclusion

Scikit-learn has been chosen to develop and run the models because of its standardised Python based model interface. WEKA was not deemed suitable because it provides no easy way to integrate with Python or NodeJS. TensorFlow would be a suitable choice but is less appropriate than scikit-learn due to the complexity it introduces.

2.3.2 Python Dependency Management

This section will outline and compare the potential choices for dependency management.

Good dependency management allows a Python project to be added to a different machine and for that project's dependencies to be installed without clashing with the existing Python dependencies of that system. Without this, the code can become tied to the developer's system-wide Python installation and prevent other Python projects from being unable to run, where they require different versions of the same package as another Python project on the system.

Conda

Conda is a high-level environment and package dependency manager which works with Python and other languages. The appeal of Conda is the simplicity of the user experience. Conda manages every aspect of environment and dependency management. This, however, comes at the cost of the complexity of the tool itself. Portability and collaboration is a concern with Conda, as to run it on a virtual machine or another developer's machine, this machine will also need Conda to be installed, which may clash with existing dependency management solutions (Conda, 2023).

Virtualenv

With virtualenv, isolated Python environments can be created and stored in a directory of the project (virtualenv, 2023).

Used in conjunction with the Python package manager, pip, dependencies can be installed to this isolated environment and their version tracked using the pip freeze command. The output of the pip freeze command is conventionally stored in a file named requirements.txt. The requirements.txt file can then be included with the source code and used to install the specific dependencies required for a project.

pipenv

The aim of pipenv is to integrate virtualenv and pip, making it a higher-level alternative to using virtualenv with pip. Practically, the benefit of pipenv is the use of a package.lock file, which allows dependencies to be locked to a specific version, providing a security benefit. This feature however does add complexity and maintainability requirements (Reitz, 2023).

Conclusion

Ultimately virtualenv was chosen for this project due to it being the most simple and lightweight solution, reducing the number of potential failure points when porting to other systems or coming back to this project in the future.

2.4 Model Type

The initial steps in model selection are rudimentary. Machine learning algorithms can be categorised as, supervised learning, semisupervised learning, unsupervised learning and reinforcement learning. Because our data set consists of structured and labelled data the appropriate choice is supervised learning, see Géron, 2017 page 24.

Predictions by the existing iArsenic models classify data points as safe, polluted or highlyPolluted. Models produced in this project therefore must produce a comparable output. Predictions from a regression model will not be as comparable to predictions from an iArsenic model as a regression model will output a continuous value. While this would produce valid predictions, it would not be applicable in this project because this estimate will have to be converted to a classification. Therefore the models considered for this project are classification based.

Rich Caruana and Alexandru Niculescu-Mizil provide a comparison of a number of supervised learning algorithms on a range of data sets and performance measures (Caruana, 2006). Based on Caruana's paper and Géron, 2017 chapter 3, the following model types have been shortlisted for their good performance across a diverse range of dataset

types:

- k-nearest neighbors classifier
- random forest classifier
- multilayer perceptron feed forward neural network classifier

2.4.1 Validating Models & Minimizing Dataset Bias

Bias is always a factor in datasets, some types of bias cannot be measured or controlled, such as sampling bias focusing on collecting data points in and around a city's capital with less focus on rural regions. This may be the case in the dataset used in this project, but further scrutinizing bias is outside the scope of this project.

The impact of ordering and sampling bias can be minimized however, this section will explore the ways this can be done.

Hold Out Validation

In hold-out validation, the dataset is split into two, a test set and a train set. The train set is used to train a model, then the model is presented with new data to make predictions for.

This proves the model can generalize to new data, validating its performance. This method is susceptible to ordering and sample bias, however, if for example the dataset is ordered alphabetically.

K-folds Cross Validation

K-folds cross-validation involves splitting the dataset into "k" splits, where k is the number of splits. Each split is then used in turn to be the test dataset and the remaining splits are used for training. Where k equals 5, a model will be trained and tested 5 times. The difference in performance can then be examined between the splits to measure the impact of selection bias (Géron, 2017 page 99).

Shuffle Split Validation

Shuffle Split validation eliminates ordering bias by randomly selecting a portion of the dataset to be used as the test split. This does however introduce random bias.

Leave One Out Validation

In leave one out validation is similar to k-folds cross-validation, but the number of folds is set to the number of data points. The result is that the model is trained once for each data point and the left out datapoint is used for the model to generate a prediction. Whether this prediction is correct or incorrect for each data point can be used to validate the performance of the model.

This method is good for models with limited data as it minimises statistical power lost by not training the model on the data in the test set. Because it requires the model to be trained and evaluated for every data point, however, it becomes less feasible for models with larger datasets and longer training and evaluation times.

Conclusion

A hybrid of shuffle split k-folds cross-validation has been used in this project. The dataset is first shuffled to eliminate ordering bias then k-folds cross-validation is used to split the data into 5 subsets, allowing us to measure sampling bias.

Leave one out validation is not viable for this project as the dataset contains over 800,000 data points. For the iArsenic model in production, model5, one model instance is approximately 11 megabytes and takes approximately 15 minutes to generate. Therefore this method would take approximately 23 years and require 8.8 terabytes to generate and store the models.

The code to generate the k-folds shuffled split data can be found at https://github.com/JavaRip/UoP-SoftEng-Dissertation/blob/main/utis/src_to_test_train.py#L40-L53.

2.4.2 Evaluation Methods

The evaluation method must be applicable to the new models and the existing iArsenic models because the iArsenic models are classification models this section will outline potential classification evaluation metrics and their strengths and weaknesses.

Accuracy / Misclassification Rate

The accuracy of a classification model measures how often the model predicts correctly, usually as a percentage. The misclassification rate is the inverse of the accuracy.

Accuracy as an evaluation metric has disadvantages. If a dataset has an 80% positive rate, a model which always predicts true will have an 80% accuracy.

$$\text{Accuracy: } \left(\frac{\text{true_positives} + \text{true_negatives}}{\text{total_predictions}} \right)$$

Specificity

The specificity measures the percentage of correctly classified negative data points; the true negative rate (Géron, 2017 page 127).

$$\text{Specificity: } \left(\frac{\text{true_negative}}{\text{true_negative} + \text{false_positive}} \right)$$

Recall / Sensitivity

The recall metric measures the percentage of correctly classified positive data points; the true positive rate (Géron, 2017 page 122).

Sensitivity: $(\frac{true_positive}{true_positive+false_negative})$

Precision

The precision metric measures the percentage of positive predictions which are truly positive (Géron, 2017 page 121).

Precision: $(\frac{true_positive}{true_positive+false_positive})$

f1

The f1 metric measures the performance of a model by using the precision and sensitivity. Unlike accuracy, this metric considers null accuracy by decreasing with incorrectly classified positive predictions by considering precision. F1 however is unaffected by the true negative rate (Géron, 2017 page 121).

f1: $(\frac{precision*sensitivity}{precision+sensitivity})$

Confusion Matrix

A confusion matrix displays the true positive rate, the true negative rate, the false positive rate and the true negative rate (Géron, 2017 page 120). While a confusion matrix can be generated for multiple classifications as well as binary classifications, the additional complexity this introduces has not been considered beneficial in achieving the desired outcomes.

		Actual	
		(+)	(-)
Predicted	(+)	Sensitivity	False Pos
	(-)	False Neg	Specificity

Receiver Operating Curve (ROC) / Area Under Curve (AUC)

The ROC plots the true positive rate over the false positive rate at different classification thresholds. This generates a curve because when the threshold is 0, every prediction is positive, leading to a true positive rate of 1 and a false positive rate of 1. At the other extreme, when the threshold is set to 1, the false positive rate will be 0 and the true positive rate will also be 0.

A model which predicts randomly will produce an AUC value on an ROC plot of 0.5, and a perfect model which predicts correctly at every threshold would produce an AUC value

of 1 (Géron, 2017 page 127). An AUC of 0 would be equivalent to 1 as the classification could be inverted.

Conclusion

The primary metric used to evaluate models in this project is the f1. Confusion matrices will also be used when providing further detail.

The selected metric must be applicable to both the iArsenic models and the machine learning based models. Because the iArsenic models do not provide a feature to change the classification threshold, a ROC cannot be produced.

2.5 Conclusion

This chapter introduced existing prediction techniques used in related fields, including iArsenic which this project builds from, contextualised the application of these models with the black box problem, explored the possible tools available to implement and evaluate new models and explored which model types have been considered for development.

The following chapters include concepts detailed in this literature review, this chapter can be referred back to while reading further chapters.

The next chapter, Methodology, will outline exactly how these concepts have been researched and implemented throughout the development of this project.

Chapter 3

Methodology

This chapter will chronologically outline the steps planned and taken in the progression of the project.

The next chapter, Experiment Design, will detail how this methodology has been implemented.

At the start of the project, a high-level representation of the research and development steps required to complete the project was created as a Gantt chart. These steps were allocated a time frame and a sprint was planned at the start of each time frame.

This workflow is best described as a modified scrum methodology with a variable and flexible sprint window for each objective. Adherence to agile principles has guided decision-making when multiple options were available.

While each sprint was different, sprints broadly involved defining a clear goal and pursuing it until the end of the sprint window. While the sprint windows were not strict cut-offs, the primary focus of my time would change upon the start of a new sprint window.

The steps detailed in this chapter deviate from the initial plan specified in the Gantt chart. The changing project plan is inline with agile principles, embracing changing requirements and changing the plan as deemed necessary during reflection between sprints.

3.1 Research Potential Projects

The first step in this project was to identify potential project ideas. This was done with brainstorming, both by thinking about the project and talking about it with others.

The project ideas broadly fell into two categories, engineering, where the objective is to produce an application, and research, where the objective is to provide insight or an answer to an unknown.

A personal goal for me in this project was to develop my skills in machine learning and

artificial intelligence. Evaluating the iArsenic models and producing alternative models with machine learning appealed to me as the most exciting project available. Because the project already contained a wealth of existing work, the objectives of this dissertation could build off this work instead of starting from scratch. This enabled higher level and higher impact objectives to be set.

As Dr Mohammed Hoque agreed to be the client to this project I was able to commit to this project, confident in knowing it had purpose and direction from the client.

3.2 Continuously Identifying Client Requirements

In line with agile principles, the priority of this project was to satisfy the clients requirements. Therefore regular communications with Dr Hoque were prioritised in the form of emails and video calls where I would provide progress updates and Dr Hoque would provide feedback and direction.

3.2.1 Clarifying the Project Objectives

Originally the vision of the project was to develop machine learning based models and compare their performance to the existing iArsenic models.

This initial vision is however does not define several key aspects of the project. How does one evaluate the performance of an iArsenic model? How does this compare to an evaluation of a machine learning based model? What is the purpose of developing new models and comparing them to the existing ones?

Ultimately the desired outcomes were clearly defined as the following:

Desired Outcomes

1. To evaluate the existing models
2. To compare the performance of the existing models with typical machine-learning-based models

Desired Outcome 1 was requested from the client. Evaluating the models informs the client of how each iteration of the models has impacted their performance and, because this model is deployed, verifies the quality of the predictions being made in deployment. While the existing models were developed with expert knowledge, prior to the evaluations conducted in this project the performance of the iArsenic models was unknown.

The purpose of desired outcome 2 builds from the assertion made in Fleming et al., 2021 that machine learning should be part of Earth and Environmental Science. iArsenic provided an opportunity for a case study where expert system models can be compared to supervised learning models.

3.3 Research Applied Machine Learning Methodologies

To design implement and evaluate machine learning models for iArsenic, I first had to research machine learning to understand what specific areas in machine learning could be applied to this dataset to achieve the desired outcomes of the project and then had to develop these models in code to gain the practical experience required.

Chapter 2 of Géron, 2017, Hands-On Machine Learning with Scikit-Learn and Tensor-Flow, was the primary resource for breaking through the barrier of theoretical knowledge, gained from written sources such as Caruana, 2006 which compares several supervised machine learning models in different applications, to actually producing my own implementations of these models.

3.4 Understanding the iArsenic Code

The iArsenic code provided the following key functions:

1. Producing a standard dataset
2. Generating iArsenic models from a dataset
3. Generating predictions from these generated models

To learn how to achieve this functionality, I would read the iArsenic code in GitHub and experiment with it as an npm package.

3.4.1 Producing a standard dataset

The iArsenic script `csv-to-json.js` makes exporting a standardised dataset trivial.

This function is implemented in the project via an package script:

```
1 "scripts": {  
2   "load-src-data": "mkdir well_data ; node  
  → node_modules/preprocessing/preprocessing/cli/csv-to-json.js -p  
  → node_modules/preprocessing/data/*.csv -o  
  → well_data/src_data.json",
```

Figure 3.1: Package.json snippet showing script to load data from iArsenic

3.4.2 Generating iArsenic models from a dataset

To verify that a model can generalise to new cases, it must be trained on a subset of the data and tested on a separate subset of data which it has not been previously exposed to, see 2.4.1 on page 17. In the iArsenic web application, the models are trained on the entire source dataset. See 3.2 on page 24 for a flowchart showing this implementation.

Generating the iArsenic models from a training subset of the data was done by passing 4 of the 5 subsets of the data to the iArsenic model generator. This was done for each of the

5 subsets of the data for each of the 3 iArsenic models. See 3.3 on 25 to see the code that achieved this.

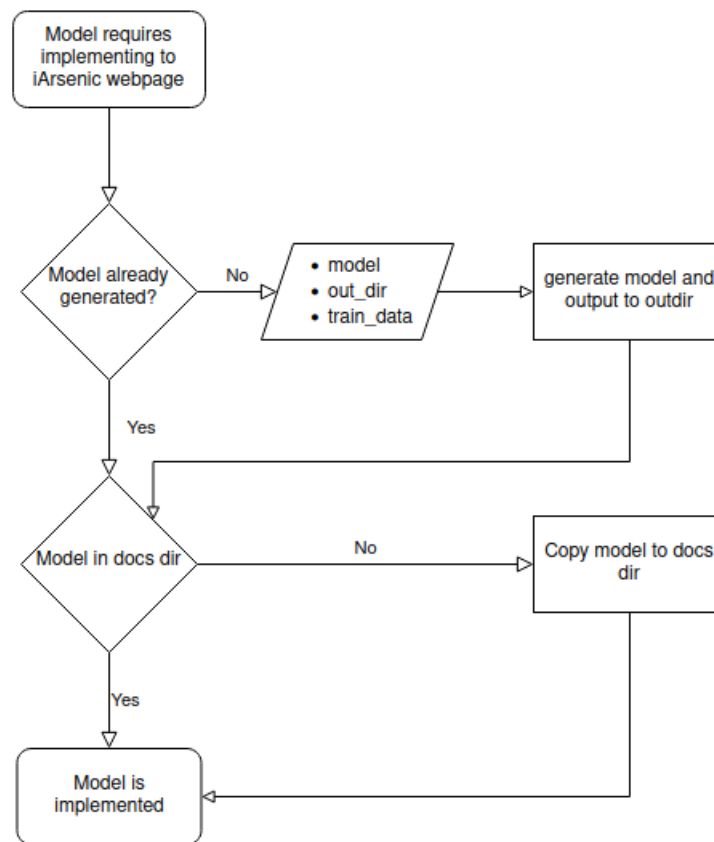


Figure 3.2: Flowchart showing process to implement model in iArsenic webapp

The code that generates the models can be found at: <https://github.com/portsoc/iArsenic/blob/master/preprocessing/cli/produce-aggregate-data-files.js>

[//github.com/portsoc/iArsenic/blob/master/preprocessing/cli/produce-aggregate-data-files.js](https://github.com/portsoc/iArsenic/blob/master/preprocessing/cli/produce-aggregate-data-files.js)

3.4.3 Generating predictions from generated iArsenic models

The web application index.html file shows a script.js file is used, in this file a global function called produceEstimate is called from the estimator.js script. See 3.4 on 25 to see this function call.

3.5 Developing New Models

To familiarise myself with implementing machine learning models, I initially followed the tutorial in Chapter 2 of Géron, 2017's book Hands-On Machine Learning with Scikit-Learn. I then implemented 4 separate supervised learning models on 25 datasets from the internet, see 3.5. This was to develop my skills, in line with the agile principle continuous attention to technical excellence.


```

1 def gen_model(train_src, out_dir, model):
2     if not os.path.exists(out_dir):
3         os.mkdir(os.path.join(out_dir))
4
5     cmd = [
6         'node',
7         'node_modules/preprocessing/preprocessing/
8         ↪ cli/produce-aggregate-data-files.js',
9         '-m',
10        model,
11        '-o',
12        out_dir,
13        '-p',
14        train_src,
15        'node_modules/preprocessing/data/mouza-names.csv',
16    ]
17    stdout = check_output(cmd).
18    decode(sys.stdout.encoding).
19    replace('\n', '')
20
21    print(stdout)

```

Figure 3.3: Python code used to programmatically generate iArsenic models. In iArsenic this is done via the command line
The file containing this code can be found at: https://github.com/JavaRip/UoP-SoftEng-Dissertation/blob/main/models/model_utils/gen_ia_model.py#L9-L25

```

1     const estimate = produceEstimate(aggregateData, inputs.division,
2     ↪ inputs.district,
3     inputs.upazila, inputs.union, inputs.mouza, inputs.depth,
4     ↪ inputs.colour, inputs.utensil, inputs.flooding);

```

Figure 3.4: Code used in iArsenic to produce estimate from a model in the webapp
The file containing this code can be found at: <https://github.com/portsoc/iArsenic/blob/master/docs/script.js#L340-L341>




Figure 3.5: Summary of results of supervised model implementation practice

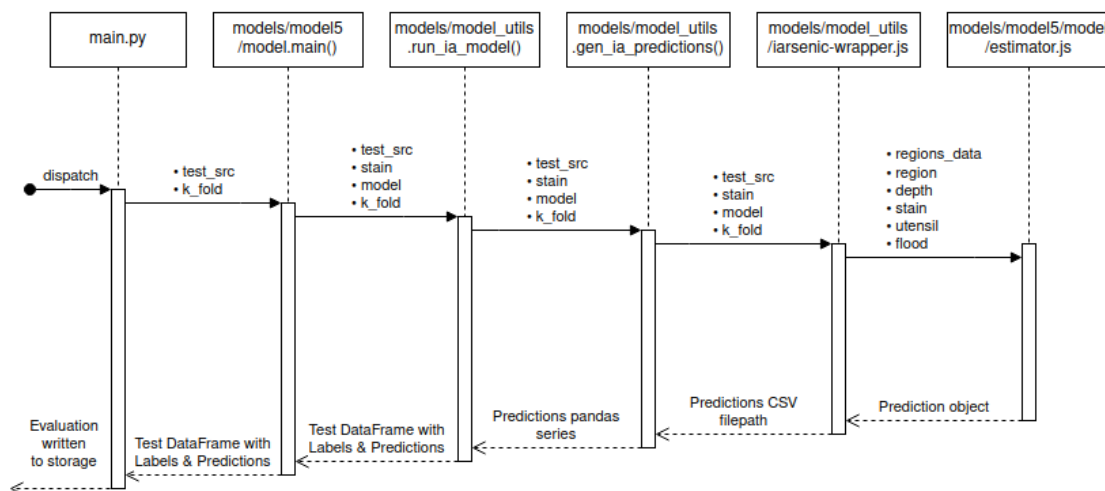


Figure 3.6: Sequence diagram showing Python to NodeJS variable bridge over storage

3.6 Integrating iArsenic and scikit-learn

The prototype implementations of both iArsenic and machine learning based models allowed each model to take input data and produce estimations individually.

Because it takes multiple days to run every experiment required, it is not practical to have a human run each model individually. The time between an experiment finishing and the next starting will depend on when the human gets around to it. By making the models run programmatically this is not an issue. It also makes the project to be more portable, allowing it to be run on a remote server or high performance computer.

To run both NodeJS based iArsenic models and Python based scikit-learn models an integration layer between Python and NodeJS had to be created. Python was chosen as the main program to run the NodeJS code from as a personal preference. Using NodeJS would also have been a reasonable choice. See 3.6 on page 27 to see the flow of data from the main.py script to the iArsenic models to produce estimates for a test dataset.

3.7 Main Project Structure

The project architecture has a top-level main file, which when run will produce an evaluation for each model. See 3.7 on page 28 to see the file structure of the project.

This structure combined with the requirement for the models to be run programmatically, facilitates continuous delivery of new models. Conforming to the agile principle of delivering working software frequently.

```
|-- evaluation_data/      # model performance metrics
|-- experiments/          # ad-hoc experiments
|-- geodata/              # uncompressed geodata
|-- geodata.zip
|-- main.py               # prediction & evaluation generation entry point
|-- models/              # all models to be run go in here
    |-- model{number}/    # model entry point & model dependencies
        |-- model.py      # takes test CSV & k-fold, returns predictions
        |-- model_utils/
|-- package.json
|-- package-lock.json
|-- prediction_data/      # test CSVs with predictions & labels by model
|-- requirements.txt
|-- utils/
|-- well_data/            # iArsenic standard dataset
```

Figure 3.7: Project file structure

3.8 Conclusion

The distinction between leadership and production in Covey, 2013 on pages 98 to 99 explains the difference between this chapter and the next.

This chapter outlines how the work is going to be done, including specifying the purpose of the project and developing the knowledge and software architecture required to produce the experiments required to achieve the desired outcomes, the next chapter Experiment Design, outlines how this plan will be implemented.

This chapter described the leadership and management, the next chapter describes the production.

Chapter 4

Experiment Design

This chapter defines how the desired outcomes defined in the Methodology chapter will practically be met. This requires describing the available code and data, how it will integrate with new code developed in this project and what the new code being developed is.

This chapter also details how the existing work in iArsenic is utilized.

4.1 Well Data in iArsenic

The data used by iArsenic can be found in the CSV in the data/ directory of iArsenic (note that this does not include CSV files in sub-directories of the data/ directory). These files have been selected on the following criteria for aggregating:

- two CSV files do not include the same datapoint
- the CSV contains standard features required by the iArsenic models

Before preprocessing, these files contain 1,144,586 rows (including headers) combined.

The primary preprocessing function provided by iArsenic is applying name corrections. This is the process of changing region names such that a region with inconsistent naming in source data files, has the same name in an aggregate dataset. These corrections allow the well data and the geodata to be aggregated.

After preprocessing, the iArsenic data contains 868,678 rows (excluding headers). 475,182 data points are negative (safe/equal to or below 10µg/l), and 393,496 are positive (polluted / above 10µg/l), the null positive rate (the accuracy when always predicting false) is 55%. The name corrected dataset output by iArsenic will be referred to as the iArsenic Standard Dataset.

See figure 4.1 on page 30 for a choropleth of datapoints per District.

For a more detailed summary of the data, see table 4.1 on page 31 to see the region sizes and datapoints per region.

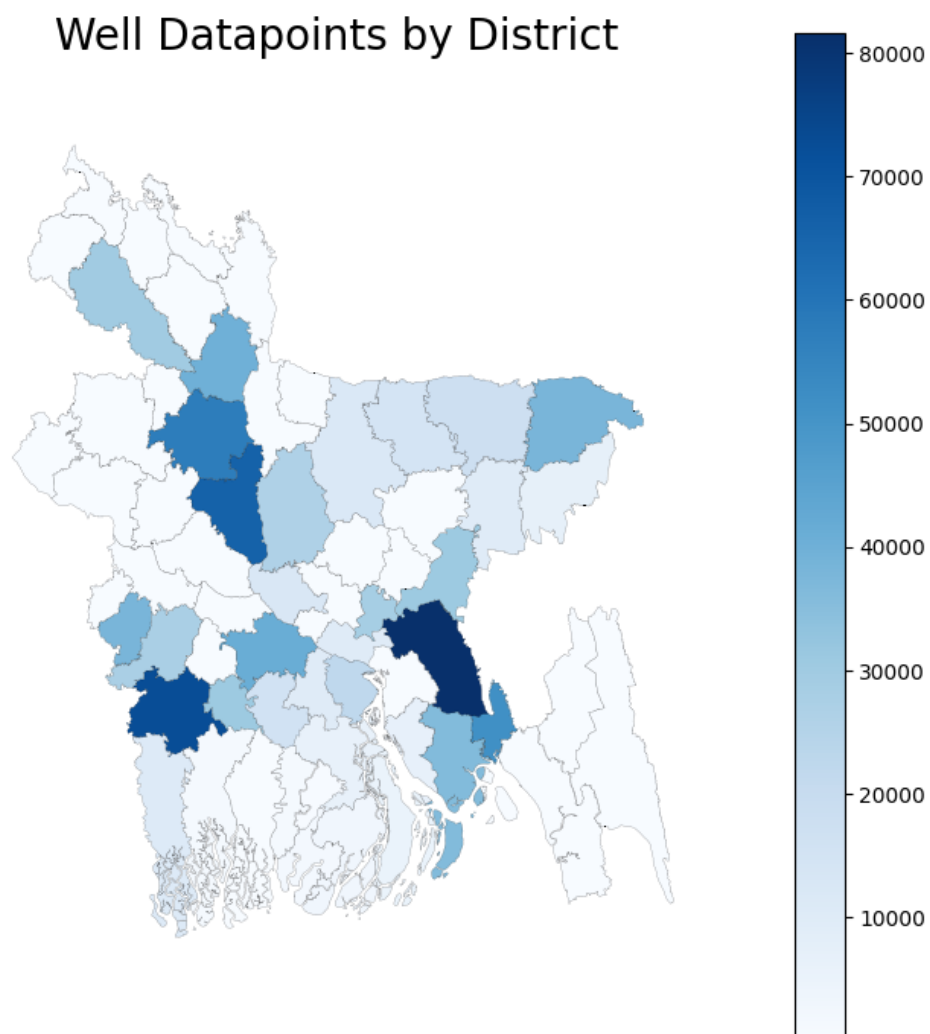


Figure 4.1: Distribution of Well Data by District in iArsenic Standard Dataset

Name	Mean Area km^2	Area σ	Mean Data	Data σ	Unique Values
Division	17,482	6,928	108,584	72,085	8
District	2,185	1,080	13,788	20,102	63
Upazila	257	185	1,952	5,683	445
Union	27	41	306	786	2,838
Mouza	2.4	8.4	90	191	9,550

Table 4.1: Region sizes, datapoints per region & count

The attributes included in the iArsenic standard dataset file are the following:

- Division
- District
- Upazila
- Union
- Mouza
- Depth
- Arsenic

Excluding Depth and Arsenic, attributes are region names with corresponding geographic data.

See figure 4.2 on page 32 for a labelled map of Bangladesh divisions.

Depth

The Depth column refers to the depth of a well in meters.

The deepest well in the dataset is 445 meters, the lowest is 0, the is mean 40 meters and the standard deviation is 52 meters. See figure 4.3 on page 33 for a choropleth of mean depth per Upazila and 4.4 on page 33 so see the relationship between mean depth and mean arsenic by division.

Arsenic

The Arsenic column refers to the arsenic concentration in water from a well datapoint in micro-grams per litre ($\mu g/l$).

The highest concentration of arsenic in the iArsenic standard dataset is $4,000\mu g/l$ ($4ng/l$), the lowest concentration of arsenic is $0\mu g/l$, the mean is $41\mu g/l$ and the standard deviation is $74\mu g/l$. See figure 4.5 on page 34 for a choropleth of mean depth per Upazila.

Divisions of Bangladesh

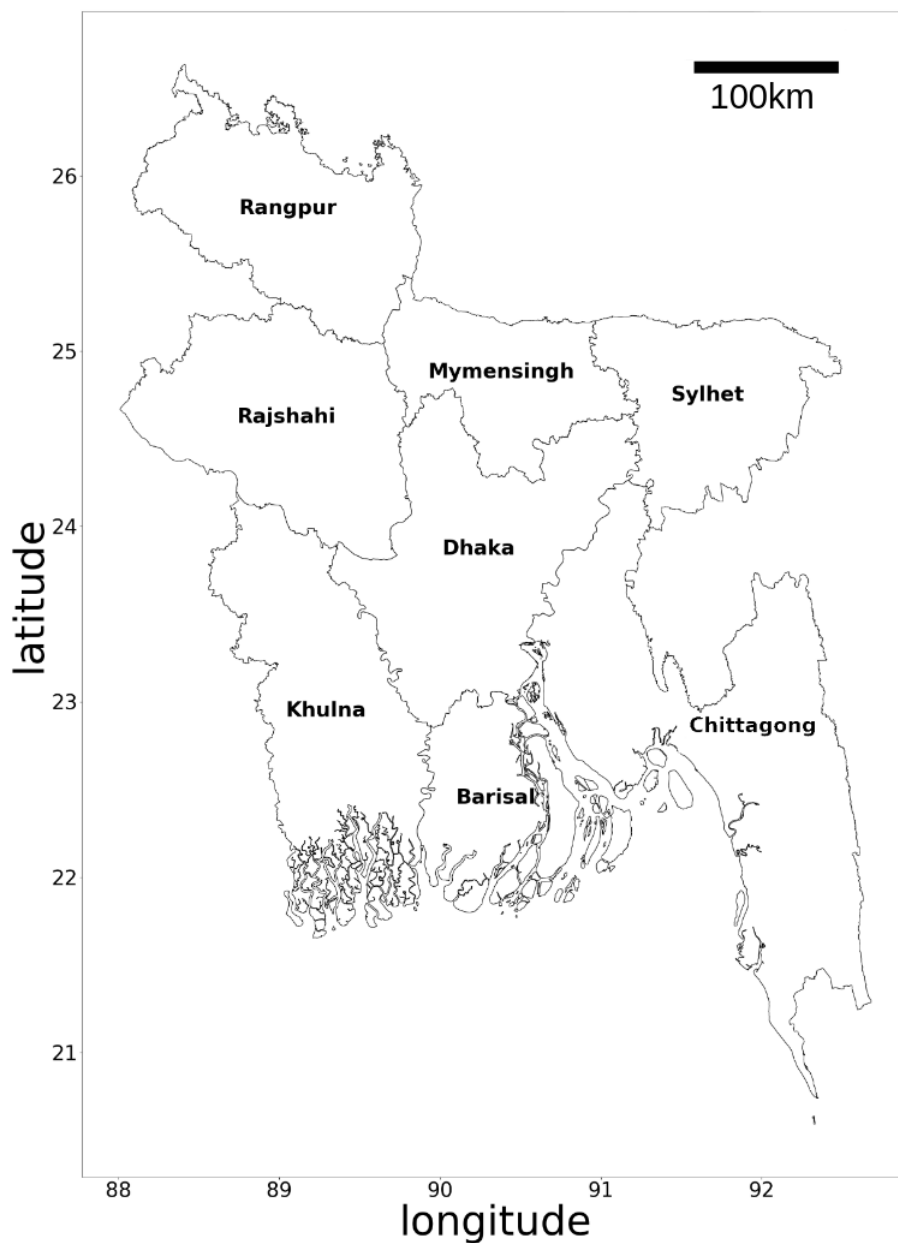
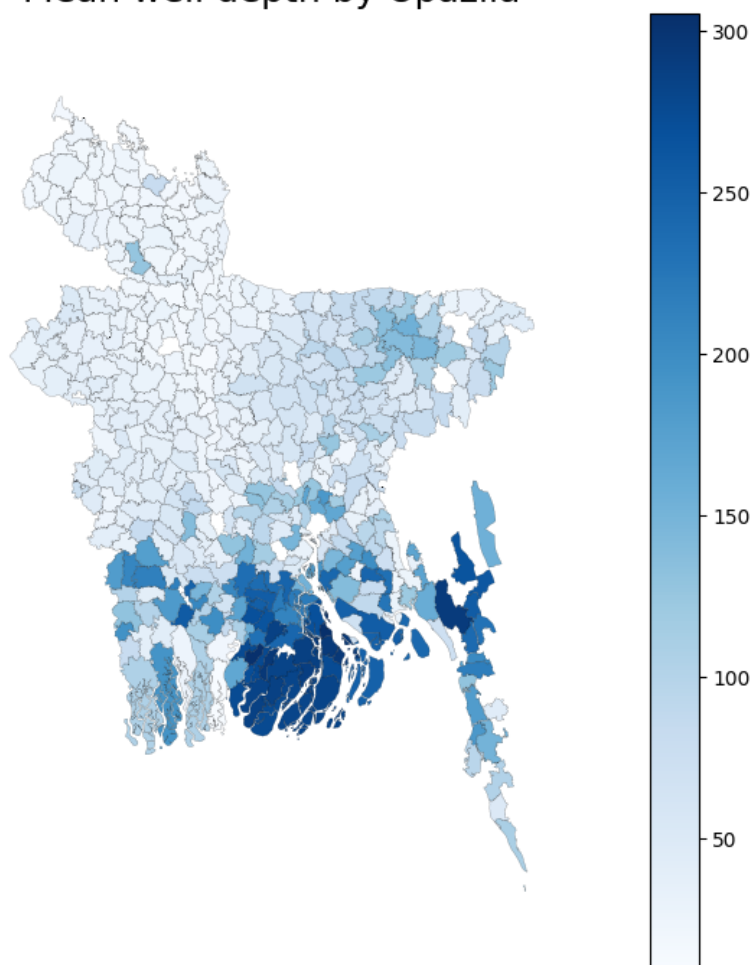



Figure 4.2: The Divisions of Bangladesh mapped and labelled

Mean well depth by Upazila

**Figure 4.3:** Average (mean) Well Depth by Upazila

Mean Depth vs Arsenic per Division

**Figure 4.4:** Chart showing the relationship between mean well depth and arsenic pollution

Mean Well Arsenic (as) by Upazila

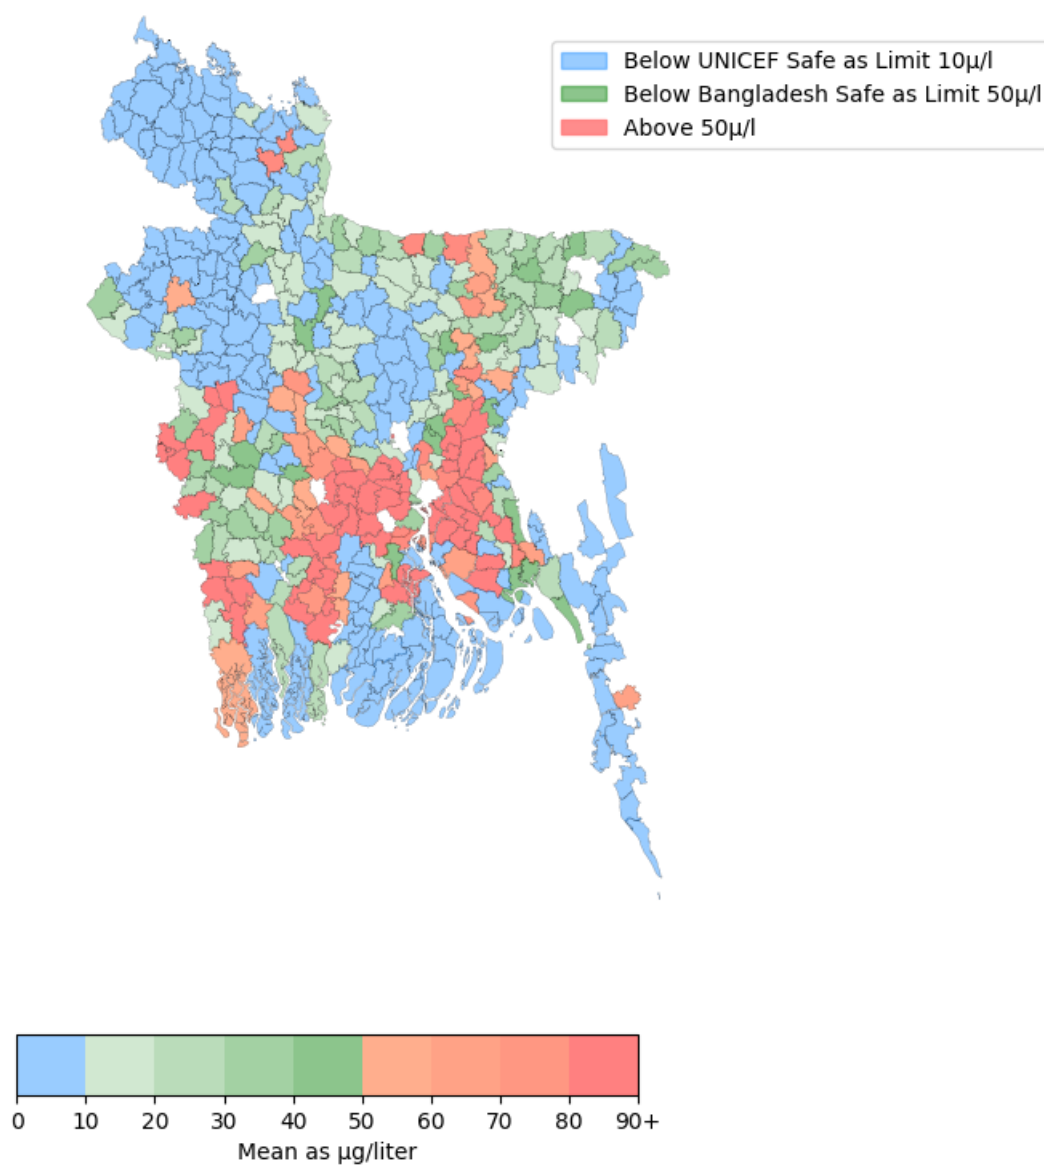


Figure 4.5: Distribution of Average (mean) Arsenic by Upazila in the iArsenic Standard Dataset

See page 79 for a red-green colorblind compatible version of this diagram

Why use this Dataset?

While new and existing iArsenic models could be trained on different data, the time and resources required to produce a dataset of this quality and suitability are substantial and outside the scope of this project.

4.1.1 Geodata

The original source of the geodata is unclear, though similar country-based polygon data can be found on the first page of Google.

The geodata has been zipped and added to the project's repository using Git Large File Storage. This is because some of the geodata are omitted from the original iArsenic repository due to it exceeding the GitHub file size limit.

Instructions are provided in iArsenic for regenerating this missing data at the following url: <https://github.com/portsoc/iArsenic/tree/master/preprocessing/geodata>

4.2 iArsenic Integration

4.2.1 Data Preparation

The iArsenic code repository is downloaded from GitHub via Node Package Manager (npm). This repository includes selected source data and several tools to process this data.

Data is extracted from iArsenic using the csv-to-json.js script. This script aggregates the source files into a single data structure and outputs it as a JSON file. Additionally, it ensures that all regions follow the same naming convention, which is essential because it prevents distinct data points in the same region from appearing to be in two different regions. This enables the data to be matched to its corresponding polygon in the geodata.

The JSON file is then converted to a pandas DataFrame in Python, shuffled to ensure there is no ordering bias, and converted into 5 separate CSV files. It is split into 5 to enable k-fold cross-validation on the dataset.

These 5 CSV files are what we use to test and train the models.

See the script that generates the 5 data subsets here: https://github.com/JavaRip/UoP-SoftEng-Dissertation/blob/main/utils/src_to_test_train.py

4.2.2 Generating the iArsenic Models

iArsenic can generate 3 separate models: model3, model4 and model5.

Originally an iArsenic model was generated on all available data and then made available on the iArsenic web application. Due to their modular design, however, the models can also be imported into NodeJS.

To evaluate the models using k-folds cross-validation, each of the 3 models is generated 5 times with 4 of the CSVs used for training and 1 for testing, thus implementing k-folds cross-validation. This requires generating a total of 15 iArsenic models.

4.2.3 Interfacing With Generated iArsenic Models

Predictions are generated from the generated model using a NodeJS script. This will be referred to as the iArsenic Model Wrapper.

The predictions are saved to a temporary CSV file and the name of this file is logged to the standard output.

The iArsenic model wrapper requires the following parameters to be passed as system arguments:

- a CSV file containing the data points to use to generate predictions
- the stain colour of the well the data point is from
- the model to generate the prediction from
- the k-fold version of the model to generate the prediction from

Initially, the iArsenic Wrapper Script would produce an estimate from one data point at a time, passed as a system argument. However, this would take days to pass the every datapoint in the test dataset. Therefore the test dataset is passed as a CSV file path, which takes less than an hour to process.

4.2.4 Imputing Stain Colour

In the original iArsenic web application, the stain colour was to be specified by a user entering parameters about a real well, requiring the user to observe and provide the staining colour of the well. The data extracted from iArsenic however did not include staining colour data.

When working with missing values there are primarily 3 options (see Géron, 2017 page 88):

1. remove the feature with missing values
2. delete all rows containing missing data
3. replace missing data with a fixed value

Because the iArsenic models will not run with the colour omitted, option 1 is not viable. Because none of the rows contain staining data, option 2 is also not viable. Therefore we must fill in the data with a fixed value, either 'Red' or 'Black'.

To determine whether the value should be set to 'Red' or 'Black', the models have been evaluated using each and the value with which the model performs best, 'Red', has been selected. Figure 4.6 on page 37 shows the accuracy of the iArsenic models with 'Black' or 'Red' used as the imputed well colour.

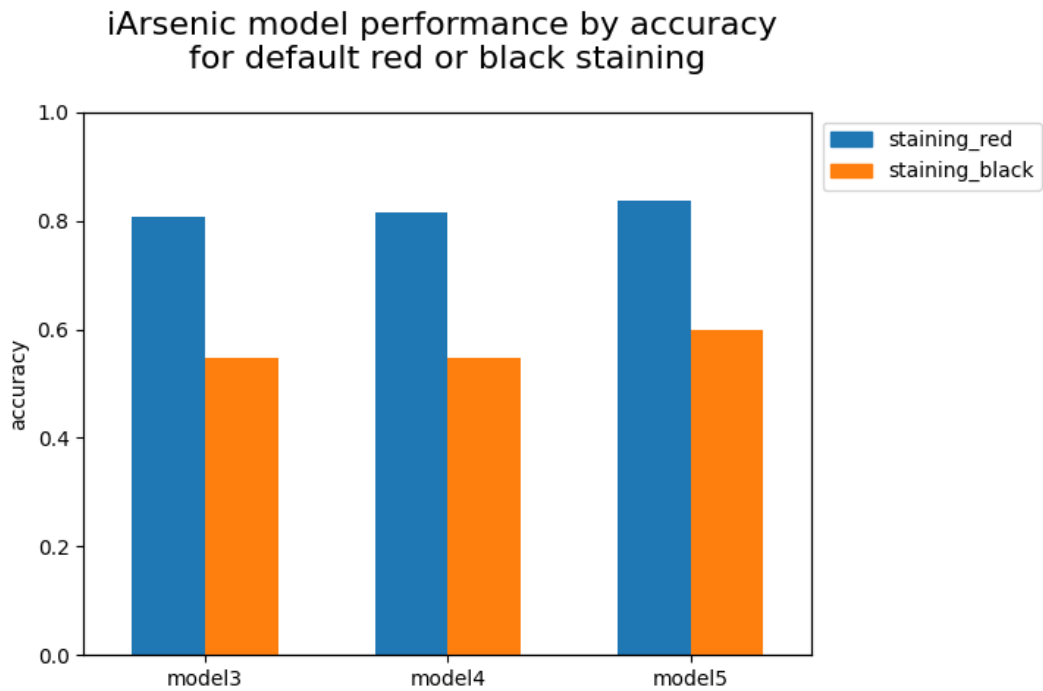


Figure 4.6: iArsenic model performance by accuracy for default red vs black staining

Considering the accuracy score alone, it appears that setting the parameter to 'Red' produces the best model performance by approximately 20%. As mentioned in **EM** on page 18, a higher accuracy score however does not guarantee that one model outperforms another.

Comparing the sensitivity and specificity of the models shows that when the staining is set to 'Black' model3 and model4 always assume the well is safe. The sensitivity, the true positive rate, becomes 0 and the specificity, the true negative rate, becomes 1. While the specificity is lower when the staining is set to 'Red', the sensitivity is much higher for each model, indicating better performance. See 4.7 on page 38 for the sensitivity and specificity plotted as a bar chart for 'Red' and 'Black' staining for comparison.

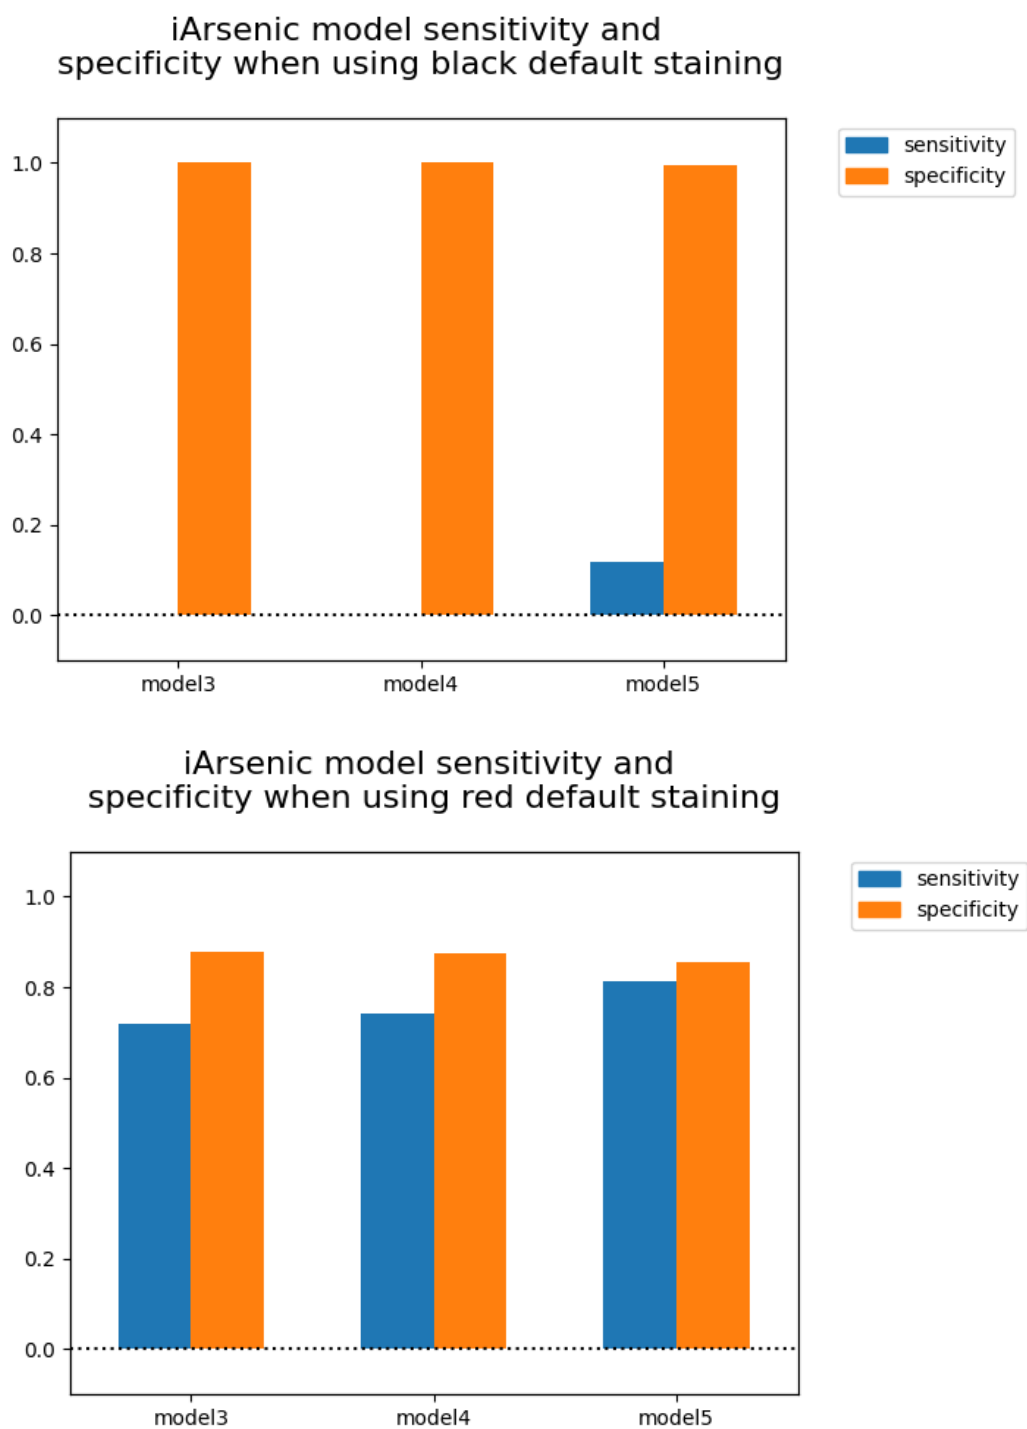


Figure 4.7: Model sensitivity and specificity stain parameter 'Black' (top) vs 'Red' (bottom)

Generating two confusion matrices for model5, one generated from 'Black' passed as a parameter and one from 'Red', reinforces the conclusion that using 'Red' as a parameter produces a better model. See 4.8 on page 39 to see a confusion matrix for each staining parameter for comparison of evaluations.

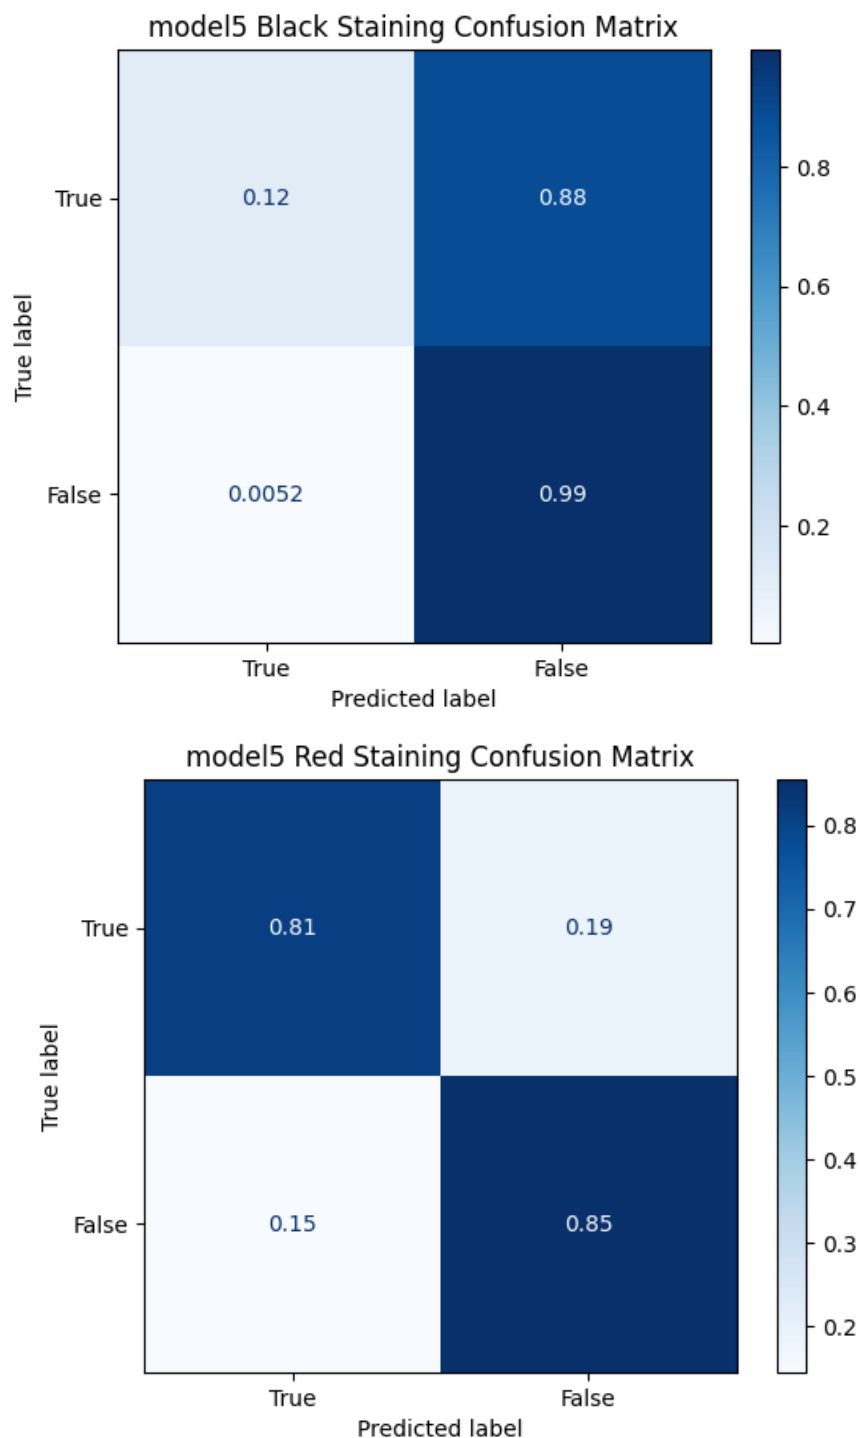


Figure 4.8: Confusion Matrix model5 'Black' (top) vs 'Red' (bottom) Staining

Quantitative Approach to Selecting Staining Parameter

Observing the evaluation metrics of the iArsenic models and comparing them when producing predictions with either 'Red' or 'Black' as the staining parameter has indicated that the models perform better with 'Red' passed as a parameter. This is however still arguably down to interpretation.

A quantitative comparison of the models, using a value such as area under the curve would be a beneficial addition to the evaluation. This however is not possible with the iArsenic models as there is no feature to change the classification threshold, which is required to generate a receiver operating curve and therefore calculate the area under the curve.

4.3 Creating a Common Model Interface

The machine learning models will be made in Python. These can be interfaced with directly using Python modules with parameters passed directly into function calls. This is inconsistent with the iArsenic models which are run via the iArsenic Wrapper Script using node, passing parameters with command line arguments.

A Common Module Interface has therefore been specified which standardises how the modules should be interacted with. This interface specifies that each model will have a main function which takes two parameters, the test data source and the k-fold number, and returns predictions as a pandas DataFrame.

4.3.1 Implementing the Common Model Interface

Each model has a directory in the top-level models directory. Each model directory contains a model.py file which exports a Python module which can generate predictions based on the Common Module Interface, in addition to any other resources required by the model. See 3.7 on page 28 for an illustration of this file structure.

Because it is a NodeJS script, the iArsenic Model Wrapper cannot accept variables passed by Python. Therefore a Python to NodeJS Bridge has been created. The Python to NodeJS Model Bridge works by running the iArsenic Model Wrapper from Python in a subprocess managed by Python from the main.py script.

4.4 Evaluating the Models

The predictions returned from the common model interface are used for evaluation. The following evaluation metrics are generated:

- accuracy
- sensitivity
- precision

- specificity
- f1

See 2.4.2 on page 18 for a detailed explanation of these metrics.

4.4.1 Standardizing Classifications

The iArsenic models generate predictions from 1 of three classifications, safe, polluted, highlyPolluted. The scikit-learn models could be configured to output any number of classifications.

To adhere to conventional evaluation metrics the iArsenic predictions were simplified to safe or polluted, where highlyPolluted predictions are converted to polluted. This enables the metrics detailed in 2.4.2 on page 18 to be used instead of defining our own evaluation methods and metrics, which would be less comparable to other real-world models.

4.5 Creating a Main Script to Run & Evaluate Experiments

The Common Module Interface allows the models imported to the main.py script and run programmatically, without specific code defined for each model.

This facilitates maintainable code as the complexity of the code increases. The complexity of the main.py script necessitates that steps to ensure maintainability are taken.

Significant complexity has been introduced in making the models run and build concurrently. This is important however as doing this sequentially would take more time than is practical.

The iArsenic models benefit the most from being run concurrently as these are single-threaded. Running them concurrently allows more of the host machine's CPU power to be utilised by using multiple CPU cores simultaneously.

The machine learning models however are written using scikit-learn libraries, which are able to use all CPU cores available.

Further optimizing the models by, for example, running the iArsenic models concurrently and the machine learning models sequentially was not deemed necessary, as building and running all models concurrently took approximately 3.5 days, which is a practical amount of time. This run was achieved on a server computer with Ubuntu 22.01, an 8 x 1.8GHz core CPU and 96GB 1600MHz DDR3 RAM.

4.5.1 Main Script Structure

The purpose of the main.py script is to provide an entry point to execute all required steps of the project.

The main script takes the following steps:

1. extract the iArsenic & geodata
2. generate the k-folds cross-validation data split
3. build the iArsenic models
4. generate predictions from all models
5. generate an evaluation for each model's predictions

4.6 Development of New Models

In the existing work section of the literature review 2.1 on page 12, 3 model approaches were examined.

The iArsenic models implement elements of k-nearest neighbour algorithms, Winkel et al., 2008 used random forests. This makes these model types candidates for further experimentation. Multilayer feed-forward neural network classifiers are also used because of their versatility in machine learning.

The models in both Winkel et al., 2008 and Connolly et al., 2022 use latitude and longitude. Because of the availability of geodata in this project, preprocessing regions to their latitude and longitude is considered a viable technique.

The iArsenic models calculate global variables during preprocessing, including the lower quartile, upper quartile and median arsenic concentrations for any given region. This processed data can be extracted from the iArsenic model and passed to a machine-learning model.

Another preprocessing technique used is applying no preprocessing.

When initially developing a model, intuition is used to hypothesize how a model might potentially work. After being implemented in code it is optimized. Neural networks have required substantial optimizations to allow the models to be trained and produce predictions within hours instead of days or weeks, at the compromise of model performance.

The predicted suitability of these model types is reinforced by Caruana, 2006 when comparing the performance of machine learning algorithms over different types of datasets.

4.6.1 Summary of Models and Preprocessing Techniques Implemented

Types of Model

1. K-Nearest Neighbours Classifier
2. Random Forest Classifier
3. Multilayer Perceptron Feed Forward Neural Network Classifier

Types of Preprocessing

1. No preprocessing
2. Conversion of region names to latitude and longitude
3. Extracting iArsenic model preprocessed data

4.6.2 model6

Model6 is an implementation of scikit-learn's random forest classifier with the default configuration.

Because this model cannot process string values, all string values are converted to numerical values, with each new string given an identification number in ascending order.

No preprocessing is applied to the data.

4.6.3 model7

Model7 is also a default implementation of scikit-learn's random forest classifier. This model incorporates computed region attributes computed in model5.

model5 uses feature engineering to calculate the median, lower quartile and upper quartile of the arsenic in a region, for a given depth range.

The hypothesis is that this feature engineering will provide the model with information that correlates with the model prediction target for a given data point, improving the models performance.

These values are therefore incorporated into the dataset used for model7. Where the values for a region are missing from the model5 dataset, the values from the region's parent are used.

4.6.4 model8

Model8 uses a feed-forward neural network, scikit-learn's MultiLayer Perceptron Classifier (MLPClassifier) model. See the configuration of model8 in table 4.2 on page 44.

Neural Network Configuration

The first hidden layer consists of half the number of inputs, the next hidden layer one-quarter the number of inputs and the final hidden layer one-eighth. The adam algorithm is used because it is generally faster than standard gradient descent on datasets this size.

Hidden Layers	Solver	Learning Rate	Epochs	Validation
3	adam	adaptive	100	10% of train

Table 4.2: Model8 MLP configuration

Hidden Layers	Solver	Learning Rate	Epochs	Validation
2	adam	adaptive	100	10% of train

Table 4.3: Model9 MLP configuration

An adaptive learning rate is used to allow the performance to increase up to the maximum limit of 100 epochs.

Data Preprocessing & Feature Engineering

The median, upper quartile and lower quartile arsenic values are imported from the model5 processed data to the dataset used by this model.

The smallest region size, the Mouzas are one hot encoded, allowing each Mouza to have a corresponding input node.

The region columns are dropped.

The data is normalised between 0 and 1.

Model Optimization

Due to the number of Mouzas (9550), the model could not allocate enough memory to run with the source data passed directly.

Therefore the source data has been split into 8 subsets by Division (See figure 4.2 on page 32. For each Division a model is trained with the data points within that region and predictions are generated. These predictions are then added to their corresponding data points in the source dataset.

The implementation of this optimization can be seen in figure 4.9 on page 48.

4.6.5 model9

model9 is also based on scikit-learn's MLPClassifier. See the configuration of model9 in table 4.3 on page 44.

Neural Network Configuration

The first hidden layer consists of 50 nodes, the next hidden layer consists of 2 nodes. An adaptive learning rate is used to allow the performance to increase up to the maximum limit of 100 epochs.

Data Preprocessing & Feature Engineering

The source data is merged with the geodata and each datapoint is attributed with a latitude and longitude. All other feature columns are dropped.

Simplifying the dataset to just latitude and longitude enables the entire dataset to be used to train and test the model.

4.6.6 model10

model10 uses scikit-learn's of the k-nearest neighbours classification algorithm, the KNeighborsClassifier. The number of neighbours is set to 50.

Data Preprocessing & Feature Engineering

The source data is merged with the geodata and the latitude and longitude of each Mouza is attributed to each datapoint. All other feature columns are dropped.

This model will find the 50 data points with the smallest difference in terms of latitude, longitude and depth in the training set then make a prediction from the classification of these data points.

4.6.7 model11

Like model10, model11 uses scikit-learn's implementation of the k-nearest neighbours classification algorithm. The number of neighbours is set to 250.

Data Preprocessing & Feature Engineering

This model uses the data processed by model5, including the lower median and upper quartile arsenic values for a region.

This model will find the 250 data points with the smallest difference in the median, lower quartile and upper quartile values of arsenic in the training dataset and make a prediction from the most common classification of those nearby data points.

4.6.8 model12

Model12, an ensemble model, uses prediction data from all models to identify which model performs best in which region.

For any data point, the model which performs best in this region could be used to generate the predictions, utilising the best of all models.

As of the date of submission, model12 is a work in progress because of limitations of the common model interface, this will be detailed further in the Future Work chapter.

A file showing which model performed best for each k-fold can be seen here: https://github.com/JavaRip/UoP-SoftEng-Dissertation/blob/main/models/model12/best_model_by_upa.csv

Hidden Layers	Solver	Learning Rate	Epochs	Validation
2	adam	adaptive	100	20% of train

Table 4.4: Model14 MLP configuration

4.6.9 model13

Like model10 and model11, model13 uses scikit-learn's k-nearest neighbours classification algorithm, the number of neighbours is set to 5000 neighbours.

Figure 4.10 on page 49 shows how arsenic levels link to depth. This illustrates why categorizing a datapoint based on data points in another depth category will indicate a lower or higher arsenic concentration.

Data Preprocessing & Feature Engineering

Model13 trains a separate model for each depth strata. This eliminates the possibility of all nearest neighbours of a data point being from the same region at a different depth.

Comparing a neighbouring data point at a different depth is wrong because deeper wells will almost always be safer. Therefore wells in the same depth range should be used as neighbours.

In addition to being split into separate depth categories, all features are removed apart from the latitude and longitude.

4.6.10 Model14

Like model8, model14 splits the data into smaller subsets. Instead of splitting the data by the Division, it is split by the depth category. See 4.4 on page 46 to see the configuration of model14.

Data Preprocessing & Feature Engineering

The data is reduced to the latitude and longitude of each data point. Every column apart from Depth is dropped. All columns are normalised to between 0 and 1.

4.7 Conclusion

This chapter describes the landscape in which the machine learning models have been developed by detailing the practicalities of the available data and the existing iArsenic models.

This illustrates the foundation the machine learning models have been developed upon, explaining their design with regard to the requirements and opportunities presented by the existing work.

In the next chapter, Results and Evaluation, the outcome of the experiment execution is evaluated in the context of the desired outcomes defined in the Methodology in section 3.2.1 page 22.

```

1  for div in train_df['Division'].unique():
2      tr_div = train[train['Division'] == div]
3      te_div = test[test['Division'] == div]
4
5      tt_df = append_test_train(te_div, tr_div)
6
7      conv_cat_num(tt_df, 'Label')
8      tt_df = ohe_col(tt_df, ['Mouza'])
9
10     tt_df = tt_df.drop(
11         columns=[
12             'Division',
13             'District',
14             'Union',
15             'Upazila',
16         ]
17     )
18
19     cat_int_enc(tt_df)
20     tt_df = pd.DataFrame(
21         MinMaxScaler().fit_transform(tt_df),
22         columns=tt_df.columns
23     )
24
25     te_div, tr_div = split_test_train(tt_df)
26
27     train_X = tr_div.drop(['Arsenic', 'Label'], axis='columns')
28     train_y = tr_div['Label']
29     test_X = te_div.drop(['Arsenic', 'Label'], axis='columns')
30
31     num_feat = len(test_X.columns)
32
33     clf = MLPClassifier(
34         solver='adam',
35         alpha=0.0001,
36         hidden_layer_sizes=(
37             math.trunc(num_feat / 2),
38             math.trunc(num_feat / 4),
39             math.trunc(num_feat / 8)
40         ),
41         learning_rate='adaptive',
42         random_state=99,
43         max_iter=100,
44     )
45
46     clf.fit(train_X, train_y)
47
48     test.loc[test['Division'] == div, ['Prediction']] = clf.predict(test_X)

```

Figure 4.9: Snippet showing m8 optimization method, where a different model is trained for each of the 8 Divisions

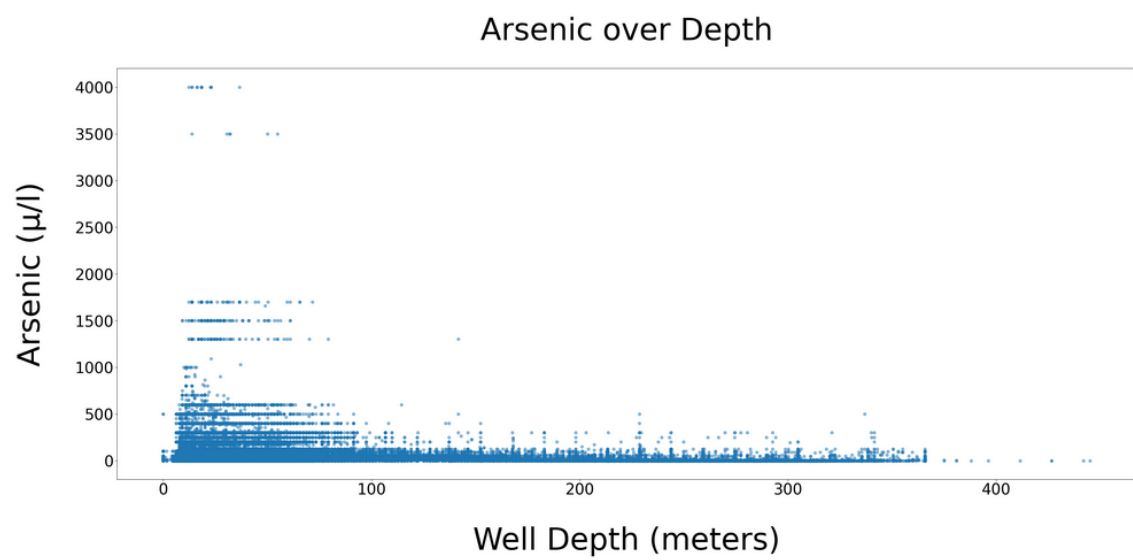


Figure 4.10: Distribution of Well Data by District in iArsenic Processed Dataset

Chapter 5

Results and Evaluation

In section 3.2.1 on page 22, the desired outcomes of the project are described. This chapter critically assesses to what extent these outcomes have been met.

5.1 Verifying Bias in Test Train Split

This section examines the limitations of the evaluation. This is done by measuring the bias in the data, using k-folds cross-validation, shuffle split and hold-out testing.

Note that not all types of bias, such as selection bias, can be measured this way and therefore the impact of these biases remains unknown.

5.1.1 Quantifying Bias

Bias in evaluation is inevitable. In section 4.2.1 on page 35 we specified that the data is shuffled to eliminate ordering bias. However, shuffling the data introduces random bias.

By using k-folds cross-validation we can generate 5 confusion matrices for each model, one for each k-fold, and then we can measure the standard deviation in the confusion matrix metrics.

Table 5.1 on page 51 shows the mean standard deviation of the true positive and false negative metric between the 5 train test splits is $4.1e-03$ and $3.39e-03$ respectively. This is equivalent to 0.41% and 0.39% standard deviation.

For the purpose of evaluating the models in this project, this deviation between splits is considered insignificant. Therefore evaluations are made on one of the k-fold test train subsets, k1.

Model	True Positive σ	True Negative σ
model3	1.75e-03	5.36e-04
model4	1.78e-03	8.51e-04
model5	1.73e-03	1.14e-03
model6	4.07e-03	3.27e-03
model7	2.15e-03	4.2e-03
model8	1.54e-03	3.32e-03
model9	2.21e-02	1.71e-02
model10	2.27e-03	1.45e-03
model11	1.95e-03	1.41e-03
model13	1.64e-03	6.36e-04
model14	2.182838e-03	6.061164e-04
mean	4.1e-03	3.39e-03

Table 5.1: Standard deviation of specificity and sensitivity across each test train split

5.2 Desired Outcome 1: Evaluating iArsenic Models

This section details how desired outcome 1 has been achieved. See the desired outcomes at 3.2.1 on page 22

Table 5.2 on page 55 shows evaluation metrics of the iArsenic and the machine learning based models, and chart 5.1 on page 54 shows the f1 scores in this table as a bar chart.

The iArsenic models, model3, model4 and model5, have been developed sequentially with each model improving on the one before. The evaluations show that the models have been improving with each iteration. The evaluation also shows that the best performing iArsenic model, model5, achieved an accuracy of 84% and an f1 of 82% and a true negative rate of 85%. See 4.8 on page 39 for a confusion matrix of model5.

Compared to the other models in the literature review, this appears to be a very high level of performance. Though further analysis is required to confirm this as an evaluation metric common to each of these models must be identified.

5.2.1 Benefit of Evaluation of iArsenic Models

Evaluating the iArsenic models validates their performance and quantifies their suitability for implementation in the iArsenic web application. This facilitates a more informed application of predictions from these models.

The evaluation shows that the performance of the models is increasing with each iteration from model3 to model5, indicating that performance has not plateaued so the development of further models may produce further improvements in performance.

If new iArsenic models are developed, the evaluation of the existing models will be able

to indicate where the existing models are underperforming, aiding the development of new models.

The benefit of this evaluation supports the notion that geoscientists should about learn machine learning and artificial intelligence, as the tools and techniques required to generate this kind of evaluation are core in the field of machine learning.

5.2.2 Desired Outcome Achievement

Prior to this project, there was no evaluation of the iArsenic models. While each iteration was more sophisticated than the last and was presumed to perform better, this was not verifiable.

Because of the evaluations done in this project, it is now verifiable that the models do perform well.

The confusion matrices produced for each model offer a detailed evaluation of the performance of each model. While other evaluation techniques such as receiver operator curves were considered, not all models were able to generate them.

Confusion matrix metrics are displayed in table 5.2 on page 55.

5.3 Desired Outcome 2: Comparing iArsenic to Machine Learning Models

Fleming et al., 2021 makes the case that because of the prevalence of artificial intelligence in earth and environmental science, earth and environmental science students should be taught machine learning and artificial intelligence skills to give geoscientists the ability to work with machine learning concepts more effectively than they currently do, usually with self taught skills.

This project provides an opportunity for a case study wherein existing expert system models, developed by the geoscientist Dr Mohammed Hoque for iArsenic, are compared with machine learning based models.

This section will analyse the benefits and limitations machine learning has brought to iArsenic throughout the development of this project.

5.3.1 iArsenic & Machine Learning Models Performance

By f1 score, the best performing iArsenic model, model5, achieves a score of 82%. The best machine learning model is model6, with an f1 score of 72%. This demonstrates that machine learning models do not inherently produce better performance than expert system models.

This indicates that the models made by a geoscientist without a background in machine learning outperformed machine learning based models. This could indicate that geoscientists may not benefit from machine learning and artificial intelligence education.

5.3.2 Desired Outcome Achievement

The development and comparison of iArsenic models to machine learning models show that machine learning models are able to produce generalised predictions. This proves that machine learning models are applicable to the iArsenic dataset in a comparable way to the existing models. It is probable that with further work machine learning models, like the ones developed in this project, could outperform the existing models.

5.4 Conclusion

Evaluating the iArsenic models and comparing them to machine learning based models has been a success.

The evaluation of the iArsenic models is a valuable contribution to the application and development of the iArsenic models and iArsenic more broadly. The tools and techniques required to generate these evaluations are rooted in machine learning and artificial intelligence. It could be argued, specifically in the case of iArsenic, that if these skills were more prevalent in geoscience, this would provide a significant benefit. This reinforces the

case made in Fleming et al., 2021.

This project also shows that machine learning is no silver bullet. While the benefit of evaluation is clear, the best performing models in this project were developed specifically from a geoscience background. Indicating that machine learning is not required by geoscientists to make predictive models. This arguably weakens the case made in Fleming et al., 2021, though with refinement it is possible that machine learning based models could perform at least as well as the expert models, especially if more sophisticated model frameworks are considered.

To conclude, while machine learning skills are valuable, increasing collaboration between geoscience and computer science may be a more practical route to the same effect as increasing machine learning education in geoscience.




Figure 5.1: Model performance by f1 score

Model	True Positive	True Negative	f1	accuracy	precision
model3	0.719387	0.880010	0.771552	0.807386	0.831874
model4	0.741016	0.876879	0.784060	0.815450	0.832413
model5	0.814367	0.854302	0.818085	0.836246	0.821837
model6	0.725523	0.763088	0.720984	0.746103	0.716502
model7	0.378534	0.623378	0.412596	0.512674	0.453395
model8	0.414230	0.587080	0.432714	0.508927	0.452924
model9	0.410958	0.753969	0.480913	0.598880	0.579569
model10	0.447443	0.706891	0.497044	0.589187	0.559013
model11	0.421738	0.579780	0.437537	0.508081	0.454565
model13	0.663838	0.503931	0.585893	0.576156	0.524328
model14	0.663940	0.503868	0.585936	0.576167	0.524335

Table 5.2: Table of evaluation metrics for all models

Chapter 6

Limitations & Future Work

6.1 Limitations

Test Train Split Aggregation

In some of the models, the test and train data are combined to normalise the data. This is to prevent data points in the test split with features larger than any data point in the test split exceeding a normalized value of 1.

This is not correct as the train set should be isolated from the model during training.

While this does negatively impact the credibility of the evaluations of the models which have used this technique, the difference made is not likely to be significant enough to change the conclusions or future work in this project.

Common Model Interface Removes Index

The common model interface has been designed to prioritise integration between the NodeJS models and the scikit-learn models. Because both of these model types can read a CSV file, the interface is restricted to passing a CSV file path.

By passing a CSV file path, the DataFrame index is not kept. This blocks model12 from splitting a DataFrame, passing these splits to other models, and then reinserting the predictions into the same row the prediction came from.

Model12 therefore cannot be completed due to a limitation of the common model interface, a solution where the index is maintained would allow models of this type to be explored.

Latitude & Longitude Preprocessing

Currently, the latitude and longitude preprocessing uses the centroids on the Unions only, not the Mouzas. This is because the function used to merge the geodata to the well data

is unable to match the geodata Mouzas and the well data Mouzas, the reason why is not clear.

With the Mouza centroids, higher resolution latitudes and longitudes could be provided to the models, allowing for more precise decision boundaries to be found.

Further Analysis of Neural Network Models

Further analyses of the neural networks will identify what their current limitations are and how they should be further developed to overcome these limitations.

Figure 6.1 on page 58 shows the loss curve of model9 with 2 hidden layers of 50 nodes then 2 nodes. This shows that the decrease in loss is plateauing at 30 iterations, suggesting that further iterations will not provide a significant performance improvement.

Figure 6.2 on page 58 shows the same neural network but with 3 layers of hidden nodes with 2000, 1000 and 100 nodes respectively. After 30 iterations, this model outperforms the model9 instance with less hidden nodes in terms of both performance and training loss. The tradeoff is more computational power being required, causing longer training and test time per iteration.

The larger model, figure 6.2, is continuing to decrease loss by iteration 30. The larger model does however show signs of overfitting as the difference between the test and the train loss is increasing. This could potentially be mitigated with a larger validation split.

This shows that more hidden nodes in the neural network produces better performance, though it has not been practical to test these models due to hardware limitations and time constraints.

6.2 Future Work

6.2.1 Further Data Integration

Flooding

Connolly et al., 2022 states that flooding is a "key driver" of groundwater arsenic levels in southeast Asia. Preliminary investigations show that flooding data does exist, but is not readily available in a format which can be integrated into this project without significant preprocessing.

This preprocessing of available flooding data to produce a generic flooding data source could improve the performance of the iArsenic models, potentially achieving new performance records, or form the basis for new machine learning based models.

River Geographic Data




Figure 6.1: Model9 loss curve 50, 2 hidden layers




Figure 6.2: Model9 loss curve 2000, 1000, 100 hidden layers

Polygon data for rivers in Bangladesh are readily available for download online. Because flooding tends to happen on riverbanks (see Connolly et al., 2022), integrating river data into the existing models could be a convenient way to see if further work on higher quality flooding data will be worthwhile.

Soil Parameters

Winkel et al., 2008 shows data from the Food and Agriculture Organization of the United Nations correlates with groundwater arsenic pollution. Integrating this into the available data could allow new models to be trained or existing models to be improved.

Elevation

Connolly et al., 2022 excluded areas with slopes greater than 0.1 degrees from their model because arsenic pollution tends to occur in flat areas. Integrating this data would enable the optimization of the machine learning models, neural network based models in particular, as sloped regions could be ignored by the model.

6.2.2 Different Model Types

Convolutional Neural Network Based Model

In section 6.1 on page 57 evidence is provided that the model9 neural network with the same design but more hidden nodes and layers perform better. But how many nodes and layers would be ideal?

There is no single rule for how many hidden nodes to use in a neural network, though in mode8, each layer uses half the number of neurons in the previous layer.

Following this rule, for 9,550 Mouzas, each with a latitude and longitude, the perfect network would have 19,100 input nodes ($9,550 * 2$) and 2 output nodes for a binary classification of 'safe' or 'polluted'. This would result in 12 hidden layers containing a total of 19,090 hidden nodes. This is an impossibly large neural network to implement.

This impossibly large design would allow the entire dataset to fit into the model without optimizing the dataset at the cost of key features.

With a convolutional neural network, these key features could be retained by passing one section of the dataset to the model at a time. The model could then identify patterns to look for within a radius of a data point's latitude and longitude instead of considering it in the context of the entire dataset. This should not impact performance as a data point is likely not dependent on far away data points.

The data is currently split by Division in model8 and depth strata in model14. These manual methods of splitting the data make assumptions about where decision boundaries should be, where a neural network should be able to identify these boundaries with training. Manually splitting the data also reduces the amount of data the model can learn from. Because a convolutional neural network would not require the data to be split into smaller datasets trained on multiple models, these disadvantages could be eliminated.

Ensemble Model

The Evaluation chapter shows that some models perform better than others, what is still not clear, however, is how often the best model outperforms the other models.

The plan for model12 is to take a training set, train and test every other model on the training set, see which model performs best for each region then pass data points in the test set to the model which performs best in that region.

A CSV of the best model for each Upazila can be see here at the following link, this data could not be used by an ensemble model however because this data has been generated from the test set: https://github.com/JavaRip/UoP-SoftEng-Dissertation/blob/main/models/model12/best_model_by_upa.csv

Deep Learning Models & Pruning

Investigation of deep learning models has been designated outside the scope of this project. Features introduced by deep learning would allow for new model paradigms to be explored and reduce the negative impact on performance inflicted by current optimization techniques.

Neural network pruning for example could allow neural network models which convert regions to their latitude and longitude to work with higher precision and more detailed region sizes (Mouzas instead of Unions for example).

Currently, an input node must exist for every point within a square which covers Bangladesh, including areas in neighbouring countries or the ocean. Implementing neural network pruning could allow these regions to be ignored by the model, instead of existing in the model to no benefit. The introduction of flooding, soil and evaluation parameters could also be aided by pruning.

6.2.3 Further Development of iArsenic Models

Classification Threshold Manipulation

The evaluation of the iArsenic models, and the machine learning models which required a common evaluation method, is limited by the inability to generate a receiver operator curve. This is because it is not possible to adjust the classification threshold of the iArsenic models. It would be possible to add this feature by changing the code in the iArsenic models.

Future Work Detailed in iArsenic Repository

The evaluation of the iArsenic models showed that they have increased in performance with each iteration. Further refinements to these models are detailed in the iArsenic repository so it is feasible that further iterations will garner further performance benefits.

6.3 Conclusion

There are four broad categories for future work, developing features to eliminate limitations, developing new models, integrating additional data and continuing development of the iArsenic models.

The most promising area of future development is the development of convolutional neural network models as this continues to expand on the research question, how do machine learning models compare to the existing iArsenic models, by introducing further machine learning paradigms.

Chapter 7

Reflection & Conclusion

7.1 Research Questions

7.1.1 What is the performance of the existing iArsenic models

Table 5.2 on page 55 shows a thorough analysis of the existing iArsenic models.

These models are deployed in the iArsenic web application and this evaluation provides a key and valuable contribution which was not available before.

This project showed that the performance of the deployed iArsenic model, model5, performs at least as well or better than other models examined in the literature review.

7.1.2 Do machine learning models have potential in the prediction of groundwater arsenic pollution?

While none of the machine learning models outperformed the iArsenic models, model6, model13 and model14 each achieved an f1 score of 58% or higher.

While model6, a random forest classifier and model13, a k-nearest neighbours classifier may be reaching their full potential in this implementation, the feed forward neural network classifiers show potential for further performance improvements with more computing power, optimization or models based on new frameworks. See 6.1 on page 57 and 6.2.2 on page 59.

Despite not outperforming the expert system models, the machine learning models achieved a performance level comparable to the iArsenic models and the other predictive models in the literature review. Therefore the answer to this research question is yes, machine learning models can predict groundwater arsenic pollution, and it is possible that with further work and potentially the use of convolutional neural networks, the performance of the iArsenic models will be exceeded.

7.2 New Models Development Strategy

When developing new models a large portion of time was spent refining the models. While refinement is important and that without refinement the machine learning models would not have performed as well, much of the time spent refining could have been spent exploring different model types.

I was personally driven by a desire to exceed the performance of the existing models, which was at no point a requirement from the client or a requirement to answer the research questions. This has led to fewer model types being explored and potentially viable model types being undiscovered.

If I was to do this project again, I would focus on developing more and a more diverse range of model types.

7.3 Achievement of Personal Goals

The selection of this project for the dissertation was influenced by my own personal motivations. This project provided an amazing opportunity to develop machine learning and artificial intelligence skills and develop experience in the field of improving access to clean drinking water. This has made working on this project an immense privilege.

In some ways, I am left feeling frustrated at the decisions I made throughout the development of this project which resulted in dead ends, wastes of time and moments where I was stuck and unable to make progress for reasons that now seem so obvious to overcome. Ultimately though this frustration stems from the wisdom I now have which I did not have at the start of this project and that represents a huge success.

Having convolutional neural networks next on the roadmap of the development of my skills simply feels unbelievable, as I feel the opportunity to have a working knowledge of cutting edge artificial intelligence and machine learning methodology is now achievable.

Developing this understanding of machine learning and working on a good, worthwhile cause has been immensely personally fulfilling and I am grateful for the opportunities this project has presented.

7.4 Conclusion

This project has been a success, answering the research questions via the desired outcomes and making a genuine contribution to an existing project, iArsenic.

There is more work to do in this project and to the benefit of iArsenic more generally, including some exciting opportunities in deep learning, with this project providing the foundation for this further work.

References

- Caruana, R. (2006). *An empirical comparison of supervised learning algorithms*. www.cs.cornell.edu
- Castelvecchi, D. (2016). The black box.
- Chakraborti, D., Rahman, M. M., Das, B., Murrill, M., Dey, S., Mukherjee, S. C., Dhar, R. K., Biswas, B. K., Chowdhury, U. K., Roy, S., Sorif, S., Selim, M., Rahman, M., & Quamruzzaman, Q. (2010). Status of groundwater arsenic contamination in bangladesh: A 14-year study report. *Water Research*, 44, 5789–5802. <https://doi.org/10.1016/j.watres.2010.06.051>
- Conda. (2023). *Conda* <https://docs.conda.io/en/latest/>.
- Connolly, C. T., Stahl, M. O., DeYoung, B. A., & Bostick, B. C. (2022). Surface flooding as a key driver of groundwater arsenic contamination in southeast asia. *Environmental Science and Technology*, 56, 928–937. <https://doi.org/10.1021/acs.est.1c05955>
- Covey, S. R. (2013). *The 7 habits of highly effective people powerful lessons in personal change*.
- Fleming, S. W., Watson, J. R., Ellenson, A., Cannon, A. J., & Vesselinov, V. C. (2021). Machine learning in earth and environmental science requires education and research policy reforms. *Nature Geoscience*, 14, 878–880. <https://doi.org/10.1038/s41561-021-00865-3>
- Géron, A. (2017). *Hands-on machine learning with scikit-learn and tensorflow: Concepts, tools, and techniques to build intelligent systems*. <http://oreilly.com/safari>
- Ghosh, G. C., Khan, M. J. H., Chakraborty, T. K., Zaman, S., Kabir, A. H., & Tanaka, H. (2020). Human health risk assessment of elevated and variable iron and manganese intake with arsenic-safe groundwater in jashore, bangladesh. *Scientific Reports*, 10. <https://doi.org/10.1038/s41598-020-62187-5>
- Google. (2023). *Tensorflow* <https://www.tensorflow.org/>.
- Guidotti, R., Monreale, A., Ruggieri, S., Turini, F., Giannotti, F., & Pedreschi, D. (2018). A survey of methods for explaining black box models. *ACM Computing Surveys*, 51. <https://doi.org/10.1145/3236009>

- Khan, A. (2014). *Groundwater studies for arsenic pollution in indus deltaic aquifers of district tando mohammad khan. sindh, pakistan: Healthenvironment hazards and migration options. department of geology, university of karachi* (Doctoral dissertation). Ph. D. dissertation, 210p.
- Khan, N. I., & Yang, H. (2014). Arsenic mitigation in bangladesh: An analysis of institutional stakeholders' opinions. *Science of the Total Environment*, 488-489, 493–504. <https://doi.org/10.1016/j.scitotenv.2013.11.007>
- Loyola-Gonzalez, O. (2019). Black-box vs. white-box: Understanding their advantages and weaknesses from a practical point of view. *IEEE Access*, 7, 154096–154113. <https://doi.org/10.1109/ACCESS.2019.2949286>
- Organization, W. H. (2003). *Arsenic in drinking-water background document for development of who guidelines for drinking-water quality*.
- Reitz, K. (2023). *Pipenv: Python dev workflow for humans* <https://pipenv.pypa.io/en/latest/>.
- scikit-learn. (2023). *Scikit-learn machine learning in python* <https://scikit-learn.org/stable/>.
- Smith, A. H., Lingas, E. O., & Rahman, M. (2000). *Contamination of drinking-water by arsenic in bangladesh: A public health emergency*.
- virtualenv. (2023). *Virtualenv* <https://virtualenv.pypa.io/en/latest/>.
- WaikatoUniversity. (2023). *Weka 3: Machine learning software in java*. <https://www.cs.waikato.ac.nz/ml/weka/index.html>
- Winkel, L., Berg, M., Amini, M., Hug, S. J., & Johnson, A. A. (2008). Predicting groundwater arsenic contamination in southeast asia from surface parameters. *Nature Geoscience*, 1, 536–542. <https://doi.org/10.1038/ngeo254>

Appendix A

Appendix

A.1 Colorblind Compatible Figures

A.2 PID



School of Computing Final Year Research Project
Project Initiation Document

Can machine learning improve geoscience
prediction models

Kane Swartz

May 4, 2023

Contents

1 Basic Details	3
2 Degree Suitability	3
3 Outline of the project environment and problem to be solved	3
3.1 Who is the client	3
3.2 What is their problem	3
3.3 Why does it need to be solved	4
4 Project aim and objectives	4
5 Project constraints	4
6 Facilities and resources	5
7 Log of risks	5
7.1 Risk 1	5
7.1.1 Description	5
7.1.2 Impact	5
7.1.3 Mitigation / Avoidance	5
7.2 Risk 2	5
7.2.1 Description	5
7.2.2 Impact	5
7.2.3 Mitigation / Avoidance	6
7.3 Risk 3	6
7.3.1 Description	6
7.3.2 Impact	6
7.3.3 Mitigation / Avoidance	6
8 Project Deliverables	6
9 Project Plan Approach	7
10 Supervisor Meetings	7
11 Legal, ethical, professional, social issues (mandatory)	7
12 Appendix A: Ethics Certificate	8
13 Appendix B: Gannt Chart	10
14 References	12

1 Basic Details

Student Name: Kane Swartz

Draft Project Title: Can machine learning improve geoscience prediction models

Course Year: Software Engineering year 3

Project Supervisor: Dr Rich Boakes

Client Organisation: University of Portsmouth School of the Environment, Geography Geosciences (SEGG)

Client Contact Name: Dr. Mohammad A Hoque

2 Degree Suitability

When I enrolled onto the Portsmouth Software Engineering course it was to develop the skills that would empower me to make impactful software. This project demonstrates a range of both software development and academic skills including problem identification, appropriate choice of software tools and the ability to do appropriate research in order to understand a specific real world problem sufficiently to specify and develop a solution. This project shows my ability to develop impactful software demonstrating that I have realised the quintessential purpose of this degree by having garnered skills to this effect.

3 Outline of the project environment and problem to be solved

3.1 Who is the client

Dr. Mohammad A Hoque is a senior lecturer in the SEGG whose ambition lies in reducing human consumption of contaminants. After procuring a BSc and MSc in Dhaka University he then went on to complete a PhD in University College London.

3.2 What is their problem

Dr. Mohammad A Hoque is a contributor to an interdisciplinary project between the SEGG and the School of Computing (SoC), iArsenic. iArsenic is a tool which predicts a data point's classification based on its features.

3.3 Why does it need to be solved

At its core the purpose of iArsenic is to contribute to the scientific effort to improve access to safe drinking water by aiding the identification of unsafe drinking water.

The need for integration between the disciplines of environmental science and machine learning has been detailed (Fleming, Watson, Ellenson, Cannon, Vesselinov, 2021, pp. 878-880).

This dissertation will work toward the goal of integrating these fields by introducing machine learning techniques to an already existing environmental science project in order to compare the accuracy of predictive models developed from an environmental science perspective to models developed from a software engineering / machine learning perspective.

4 Project aim and objectives

This is an exploratory project with the aim of comparing existing predictive models developed within the University of Portsmouth SEGG with a geoscience discipline with models developed from within a software engineering / data science discipline using machine learning techniques, which are not leveraged in the existing models.

The objective of the project will be to develop a model or multiple models, leveraging machine learning. These models will then be evaluated and their error will be compared to the existing models.

5 Project constraints

The models developed will be constrained by the objective of comparing their error to the existing models. This means the new models will need the same source data-set and produce the same format of outputs.

Time constraints and the need to produce a deliverable also come together as another fundamental constraint, meaning the research, development and review areas of the project will likely benefit from more time, almost indefinitely, requiring time to be provisioned accordingly.

6 Facilities and resources

As this project will be leveraging supervised machine learning techniques on a data-set approximately 1.3 million rows in size, a regular computer without sophisticated graphics capabilities will be suitable.

If the data-set was significantly larger or deep learning techniques were being considered, cloud computing technologies would likely be considered though this is not anticipated.

- Personal laptop
- Meetings with the client
- Occasional meeting with machine learning experts will be beneficial

7 Log of risks

7.1 Risk 1

7.1.1 Description

Time constraint

7.1.2 Impact

Almost all aspects of this project, from problem identification to algorithm research and design to execution could take more time on their own than is available for the entire project

7.1.3 Mitigation / Avoidance

The project will be split into stages and each stage will have a deadline after which the project will move on to the next stage. It is anticipated there may be temptation not to move on when the time comes, so planning ahead of time will help keep time allocation realistic.

7.2 Risk 2

7.2.1 Description

Burnout / stress related health deterioration

7.2.2 Impact

A final year dissertation is well associated with student stress. This could negatively impact not only the quality of output from myself with regards to the dissertation but could lead to more profound negative personal impact.

7.2.3 Mitigation / Avoidance

At all times I aim to prioritise my well-being and if the strain appears significant enough my strategy is to remain aware that abandoning the project for the sake of my well-being is always an option and perhaps a sensible choice as my well-being is foundational to the dissertation itself.

While I don't expect to abandon the project, maintaining this frank and honest perspective will hopefully allow me to make sensible decisions with regards to my well-being and thus my productivity too.

7.3 Risk 3

7.3.1 Description

Data loss or corruption

7.3.2 Impact

This would require work to be done again which will be time inefficient. As a time sensitive project the ramifications of this would range from minor to so severe the the project cannot be delivered.

7.3.3 Mitigation / Avoidance

Work will be stored on GitHub and google drive and backed up frequently.

8 Project Deliverables

The primary deliverable will be a written comparison and analysis of the error of all of the existing models and the models to be developed.

To produce this deliverable a machine learning model will need to be developed and will also be delivered as an artefact.

9 Project Plan Approach

The primary strategy in executing this project is to break it into distinct stages. Broadly these will be:

1. Problem research / problem identification (what should the deliverable look like)
2. Solution development (developing the solution outlined by problem research)
3. Solution evaluation (reflecting on the suitability of the solution)

Stage one is already underway and formed the basis of this PID, this stage will be concluded by the end of November 2022.

Stage two will involve developing a machine learning model and a suitable evaluation method, common to the new model and the previous models, to be concluded by the middle of February 2023 with a soft deadline to have a demonstrable machine learning based model for the demo on the 3rd February 2023. Stage three will be the write up of the project and should be concluded one week before the submission deadline on 5th May 2023 allowing for at least 1 week to review and make amendments.

10 Supervisor Meetings

Regular communication with the project supervisor will be over Discord and in person meetings subject to availability.

Having worked with Dr Rich Boakes for a number of years we have developed effective though ad-hoc communication methods such as attending a first or second year lecture to borrow his time in the last 15 minutes of the class. I have confidence that the supervisor will be available for support when required in a timely fashion on a suitably regular basis.

11 Legal, ethical, professional, social issues (mandatory)

Any interpretation of the output or results of these models has to be taken in the context of experimental or research tools. This project is not intended to provide advice or guidance on whether any source of drinking water is safe and accepts no liability for the use of these models output for this purpose or any other purpose outside of academia. Drinking water pollution is a sensitive and current topic and this project will be conducted with consideration to the human impact associated with this ongoing phenomena.

12 Appendix A: Ethics Certificate



Certificate of Ethics Review

Project title: Can machine learning improve geoscience prediction models

Name:	Kane Swartz	User ID:	941226	Application date:	21/10/2022 13:58:59	ER Number:	TETHIC-2022-104051
--------------	-------------	-----------------	--------	--------------------------	------------------------	-------------------	--------------------

You must download your referral certificate, print a copy and keep it as a record of this review.

The FEC representative(s) for the **School of Computing** is/are [Haythem Nakkas](#), [Dalin Zhou](#)

It is your responsibility to follow the University Code of Practice on Ethical Standards and any Department/School or professional guidelines in the conduct of your study including relevant guidelines regarding health and safety of researchers including the following:

- [University Policy](#)
- [Safety on Geological Fieldwork](#)

It is also your responsibility to follow University guidance on Data Protection Policy:

- [General guidance for all data protection issues](#)
- [University Data Protection Policy](#)

Which school/department do you belong to?: **School of Computing**

What is your primary role at the University?: **Undergraduate Student**

What is the name of the member of staff who is responsible for supervising your project?: **Dr Rich Boakes**

Is the study likely to involve human subjects (observation) or participants?: No

Will financial inducements (other than reasonable expenses and compensation for time) be offered to participants?: No

Are there risks of significant damage to physical and/or ecological environmental features?: No

Are there risks of significant damage to features of historical or cultural heritage (e.g. impacts of study techniques, taking of samples)?: No

Does the project involve animals in any way?: No

Could the research outputs potentially be harmful to third parties?: No

Could your research/artefact be adapted and be misused?: No

Will your project or project deliverables be relevant to defence, the military, police or other security organisations and/or in addition, could it be used by others to threaten UK security?: No

Please read and confirm that you agree with the following statements: I confirm that I have considered the implications for data collection and use, taking into consideration legal requirements (UK GDPR, Data Protection Act 2018 etc.), I confirm that I have considered the impact of this work and taken any reasonable action to mitigate potential misuse of the project outputs, I confirm that I will act ethically and honestly throughout this project

Supervisor Review

As supervisor, I will ensure that this work will be conducted in an ethical manner in line with the University Ethics Policy.

Supervisor comments:

Supervisor's Digital Signature: **rich.boakes@port.ac.uk**

Date: **21/10/2022**

13 Appendix B: Gantt Chart

- primary focus
- secondary focus
- tertiary focus
- deadline

Date	02/09/2022	07/09/2022	09/09/2022	16/09/2022	23/09/2022	30/09/2022	07/10/2022	14/10/2022	21/10/2022	28/10/2022	04/11/2022	11/11/2022	18/11/2022	25/11/2022	02/12/2022	09/12/2022	16/12/2022	23/12/2022	30/12/2022	06/01/2023	13/01/2023	20/01/2023	27/01/2023	03/02/2023	10/02/2023	17/02/2023	24/02/2023	03/03/2023	10/03/2023	17/03/2023	24/03/2023	31/03/2023	07/04/2023	14/04/2023	21/04/2023	28/04/2023	05/05/2023	12/05/2023	19/05/2023	26/05/2023	02/06/2023	09/06/2023	16/06/2023	23/06/2023	30/06/2023	07/07/2023	14/07/2023	21/07/2023	28/07/2023	04/08/2023	11/08/2023	18/08/2023	25/08/2023	01/09/2023	08/09/2023	15/09/2023	22/09/2023	29/09/2023	06/10/2023	13/10/2023	20/10/2023	27/10/2023	03/11/2023	10/11/2023	17/11/2023	24/11/2023	01/12/2023	08/12/2023	15/12/2023	22/12/2023	29/12/2023	05/01/2024	12/01/2024	19/01/2024	26/01/2024	02/02/2024	09/02/2024	16/02/2024	23/02/2024	01/03/2024	08/03/2024	15/03/2024	22/03/2024	29/03/2024	05/04/2024	12/04/2024	19/04/2024	26/04/2024	03/05/2024	10/05/2024	17/05/2024	24/05/2024	31/05/2024	06/06/2024	13/06/2024	20/06/2024	27/06/2024	04/07/2024	11/07/2024	18/07/2024	25/07/2024	01/08/2024	08/08/2024	15/08/2024	22/08/2024	29/08/2024	05/09/2024	12/09/2024	19/09/2024	26/09/2024	03/10/2024	10/10/2024	17/10/2024	24/10/2024	31/10/2024	07/11/2024	14/11/2024	21/11/2024	28/11/2024	05/12/2024	12/12/2024	19/12/2024	26/12/2024	02/01/2025	09/01/2025	16/01/2025	23/01/2025	30/01/2025	06/02/2025	13/02/2025	20/02/2025	27/02/2025	06/03/2025	13/03/2025	20/03/2025	27/03/2025	03/04/2025	10/04/2025	17/04/2025	24/04/2025	01/05/2025	08/05/2025	15/05/2025	22/05/2025	29/05/2025	05/06/2025	12/06/2025	19/06/2025	26/06/2025	03/07/2025	10/07/2025	17/07/2025	24/07/2025	31/07/2025	07/08/2025	14/08/2025	21/08/2025	28/08/2025	04/09/2025	11/09/2025	18/09/2025	25/09/2025	02/10/2025	09/10/2025	16/10/2025	23/10/2025	30/10/2025	06/11/2025	13/11/2025	20/11/2025	27/11/2025	04/12/2025	11/12/2025	18/12/2025	25/12/2025	01/01/2026	08/01/2026	15/01/2026	22/01/2026	29/01/2026	05/02/2026	12/02/2026	19/02/2026	26/02/2026	05/03/2026	12/03/2026	19/03/2026	26/03/2026	02/04/2026	09/04/2026	16/04/2026	23/04/2026	30/04/2026	07/05/2026	14/05/2026	21/05/2026	28/05/2026	04/06/2026	11/06/2026	18/06/2026	25/06/2026	02/07/2026	09/07/2026	16/07/2026	23/07/2026	30/07/2026	06/08/2026	13/08/2026	20/08/2026	27/08/2026	03/09/2026	10/09/2026	17/09/2026	24/09/2026	01/10/2026	08/10/2026	15/10/2026	22/10/2026	29/10/2026	05/11/2026	12/11/2026	19/11/2026	26/11/2026	03/12/2026	10/12/2026	17/12/2026	24/12/2026	01/01/2027	08/01/2027	15/01/2027	22/01/2027	29/01/2027	05/02/2027	12/02/2027	19/02/2027	26/02/2027	05/03/2027	12/03/2027	19/03/2027	26/03/2027	02/04/2027	09/04/2027	16/04/2027	23/04/2027	30/04/2027	07/05/2027	14/05/2027	21/05/2027	28/05/2027	04/06/2027	11/06/2027	18/06/2027	25/06/2027	02/07/2027	09/07/2027	16/07/2027	23/07/2027	30/07/2027	06/08/2027	13/08/2027	20/08/2027	27/08/2027	03/09/2027	10/09/2027	17/09/2027	24/09/2027	01/10/2027	08/10/2027	15/10/2027	22/10/2027	29/10/2027	05/11/2027	12/11/2027	19/11/2027	26/11/2027	03/12/2027	10/12/2027	17/12/2027	24/12/2027	01/01/2028	08/01/2028	15/01/2028	22/01/2028	29/01/2028	05/02/2028	12/02/2028	19/02/2028	26/02/2028	05/03/2028	12/03/2028	19/03/2028	26/03/2028	02/04/2028	09/04/2028	16/04/2028	23/04/2028	30/04/2028	07/05/2028	14/05/2028	21/05/2028	28/05/2028	04/06/2028	11/06/2028	18/06/2028	25/06/2028	02/07/2028	09/07/2028	16/07/2028	23/07/2028	30/07/2028	06/08/2028	13/08/2028	20/08/2028	27/08/2028	03/09/2028	10/09/2028	17/09/2028	24/09/2028	01/10/2028	08/10/2028	15/10/2028	22/10/2028	29/10/2028	05/11/2028	12/11/2028	19/11/2028	26/11/2028	03/12/2028	10/12/2028	17/12/2028	24/12/2028	01/01/2029	08/01/2029	15/01/2029	22/01/2029	29/01/2029	05/02/2029	12/02/2029	19/02/2029	26/02/2029	05/03/2029	12/03/2029	19/03/2029	26/03/2029	02/04/2029	09/04/2029	16/04/2029	23/04/2029	30/04/2029	07/05/2029	14/05/2029	21/05/2029	28/05/2029	04/06/2029	11/06/2029	18/06/2029	25/06/2029	02/07/2029	09/07/2029	16/07/2029	23/07/2029	30/07/2029	06/08/2029	13/08/2029	20/08/2029	27/08/2029	03/09/2029	10/09/2029	17/09/2029	24/09/2029	01/10/2029	08/10/2029	15/10/2029	22/10/2029	29/10/2029	05/11/2029	12/11/2029	19/11/2029	26/11/2029	03/12/2029	10/12/2029	17/12/2029	24/12/2029	01/01/2030	08/01/2030	15/01/2030	22/01/2030	29/01/2030	05/02/2030	12/02/2030	19/02/2030	26/02/2030	05/03/2030	12/03/2030	19/03/2030	26/03/2030	02/04/2030	09/04/2030	16/04/2030	23/04/2030	30/04/2030	07/05/2030	14/05/2030	21/05/2030	28/05/2030	04/06/2030	11/06/2030	18/06/2030	25/06/2030	02/07/2030	09/07/2030	16/07/2030	23/07/2030	30/07/2030	06/08/2030	13/08/2030	20/08/2030	27/08/2030	03/09/2030	10/09/2030	17/09/2030	24/09/2030	01/10/2030	08/10/2030	15/10/2030	22/10/2030	29/10/2030	05/11/2030	12/11/2030	19/11/2030	26/11/2030	03/12/2030	10/12/2030	17/12/2030	24/12/2030	01/01/2031	08/01/2031	15/01/2031	22/01/2031	29/01/2031	05/02/2031	12/02/2031	19/02/2031	26/02/2031	05/03/2031	12/03/2031	19/03/2031	26/03/2031	02/04/2031	09/04/2031	16/04/2031	23/04/2031	30/04/2031	07/05/2031	14/05/2031	21/05/2031	28/05/2031	04/06/2031	11/06/2031	18/06/2031	25/06/2031	02/07/2031	09/07/2031	16/07/2031	23/07/2031	30/07/2031	06/08/2031	13/08/2031	20/08/2031	27/08/2031	03/09/2031	10/09/2031	17/09/2031	24/09/2031	01/10/2031	08/10/2031	15/10/2031	22/10/2031	29/10/2031	05/11/2031	12/11/2031	19/11/2031	26/11/2031	03/12/2031	10/12/2031	17/12/2031	24/12/2031	01/01/2032	08/01/2032	15/01/2032	22/01/2032	29/01/2032	05/02/2032	12/02/2032	19/02/2032	26/02/2032	05/03/2032	12/03/2032	19/03/2032	26/03/2032	02/04/2032	09/04/2032	16/04/2032	23/04/2032	30/04/2032	07/05/2032	14/05/2032	21/05/2032	28/05/2032	04/06/2032	11/06/2032	18/06/2032	25/06/2032	02/07/2032	09/07/2032	16/07/2032	23/07/2032	30/07/2032	06/08/2032	13/08/2032	20/08/2032	27/08/2032	03/09/2032	10/09/2032	17/09/2032	24/09/2032	01/10/2032	08/10/2032	15/10/2032	22/10/2032	29/10/2032	05/11/2032	12/11/2032	19/11/2032	26/11/2032	03/12/2032	10/12/2032	17/12/2032	24/12/2032	01/01/2033	08/01/2033	15/01/2033	22/01/2033	29/01/2033	05/02/2033	12/02/2033	19/02/2033	26/02/2033	05/03/2033	12/03/2033	19/03/2033	26/03/2033	02/04/2033	09/04/2033	16/04/2033	23/04/2033	30/04/2033	07/05/2033	14/05/2033	21/05/2033	28/05/2033	04/06/2033	11/06/2033	18/06/2033	25/06/2033	02/07/2033	09/07/2033	16/07/2033	23/07/2033	30/07/2033	06/08/2033	13/08/2033	20/08/2033	27/08/2033	03/09/2033	10/09/2033	17/09/2033	24/09/2033	01/10/2033	08/10/2033	15/10/2033	22/10/2033	29/10/2033	05/11/2033	12/11/2033	19/11/2033	26/11/2033	03/12/2033	10/12/2033	17/12/2033	24/12/2033	01/01/2034	08/01/2034	15/01/2034	22/01/2034	29/01/2034	05/02/2034	12/02/2034	19/02/2034	26/02/2034	05/03/2034	12/03/2034	19/03/2034	26/03/2034	02/04/2034	09/04/2034	16/04/2034	23/04/2034	30/04/2034	07/05/2034	14/05/2034	21/05/2034	28/05/2034	04/06/2034	11/06/2034	18/06/2034	25/06/2034	02/07/2034	09/07/2034	16/07/2034	23/07/2034	30/07/2034	06/08/2034	13/08/2034	20/08/2034	27/08/2034	03/09/2034	10/09/2034	17/09/2034	24/09/2034	01/10/2034	08/10/2034	15/10/2034	22/10/2034	29/10/2034	05/11/2034	12/11/2034	19/11/2034	26/11/2034	03/12/2034	10/12/2034	17/12/2034	24/12/2034	01/01/2035	08/01/2035	15/01/2035	22/01/2035	29/01/2035	05/02/2035	12/02/2035	19/02/2035	26/02/2035	05/03/2035	12/03/2035	19/03/2035	26/03/2035	02/04/2035	09/04/2035	16/04/2035	23/04/2035	30/04/2035	07/05/2035	14/05/2035	21/05/2035	28/05/2035	04/06/2035	11/06/2035	18/06/2035	25/06/2035	02/07/2035	09/07/2035	16/07/2035	23/07/2035	30/07/2035	06/08/2035	13/08/2035	20/08/2035	27/08/2035	03/09/2035	10/09/2035	17/09/2035	24/09/2035	01/10/2035	08/10/2035	15/10/2035	22/10/2035	29/10/2035	05/11/2035	12/11/2035	19/11/2035	26/11/2035	03/12/2035	10/12/2035	17/12/2035	24/12/2035	01/01/2036	08/01/2036	15/01/2036	22/01/2036	29/01/2036	05/02/2036	12/02/2036	19/02/2036	26/02/2036	05/03/2036	12/03/2036	19/03/2036	26/03/2036	02/04/2036	09/04/2036	16/04/2036	23/04/2036	30/04/2036	07/05/2036	14/05/2036	21/05/2036	28/05/2036	04/06/2036	11/06/2036	18/06/2036	25/06/2036	02/07/2036	09/07/2036	16/07/2036	23/07/2036	30/07/2036	06/08/2036	13/08/2036	20/08/2036	27/08/2036	03/09/2036	10/09/2036	17/09/2036	24/09/2036	01/10/2036	08/10/2036	15/10/2036	22/10/2036	29/10/2036	05/11/2036	12/11/2036	19/11/2036	26/11/2036	03/12/2036	10/12/2036	17/12/2036	24/12/2036	01/01/2037	08/01/2037	15/01/2037	22/01/2037	29/01/2037	05/02/2037	12/02/2037	19/02/2037	26/02/2037	05/03/2037	12/03/2037	19/03/2037	26/03/2037	02/04/2037	09/04/2037	16/04/2037	23/04/2037	30/04/2037	07/05/2037	14/05/2037	21/05/2037	28/05/2037	04/06/2037	11/06/2037	18/06/2037	25/06/2037	02/07/2037	09/07/2037	16/07/2037	23/07/2037	30/07/2037	06/08/2037	13/08/2037	20/08/2037	27/08/2037	03/09/2037	10/09/2037	17/09/2037	24/09/2037	01/10/2037	08/10/2037	15/10/2037	22/10/2037	29/10/2037	05/11/2037	12/11/2037	19/11/2037	26/11/2037	03/12/2037	10/12/2037	17/12/2037	24/12/2037	01/01/2038	08/01/2038	15/01/2038	22/01/2038	29/01/2038	05/02/2038	12/02/2038	19/
------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	------------	-----

14 References

Fleming, S. W., Watson, J. R., Ellenson, A., Cannon, A. J., Vesselinov, V. C. (2021). Machine learning in Earth and environmental science requires education and research policy reforms. *Nature Geoscience*, 14(12), 878–880. <https://doi.org/10.1038/s41561-021-00865-3>.

Mean Well Arsenic (as) by Upazila

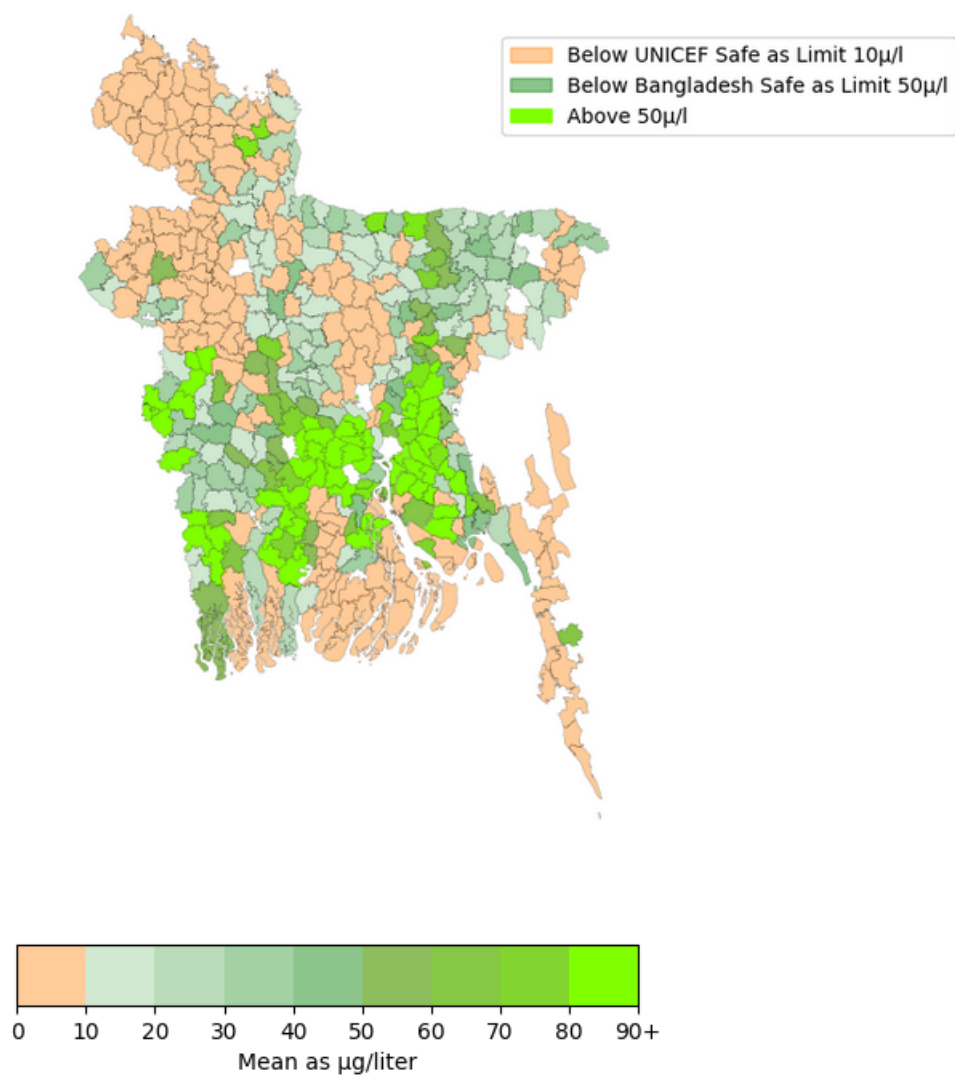


Figure A.1: Red excluded for red / green colourblind: Distribution of Average (mean) Arsenic by Upazila in the iArsenic Standard Dataset